
Best-of-Both Worlds for linear contextual bandits with paid observations

Anonymous Author

Abstract

We study the problem of linear contextual bandits with paid observations, where at each round the learner selects an action in order to minimize its loss in a given context, and can then decide to pay a fixed cost to observe the loss of any arm. Building on the Follow-the-Regularized-Leader framework with efficient estimators via Matrix Geometric Resampling, we introduce a computationally efficient Best-of-Both-Worlds (BOBW) algorithm for this problem. We show that it achieves the minimax-optimal regret of $\Theta(T^{2/3})$ in adversarial settings, while guaranteeing polylogarithmic regret in (corrupted) stochastic regimes. Our approach builds on the framework from Tsuchiya and Ito [2024] to design BOBW algorithms for “hard problem”, using analysis techniques tailored for the setting that we consider.

1 INTRODUCTION

Multi-armed bandits (MAB) have emerged as one of the most popular models for sequential decision-making under uncertainty [Lattimore and Szepesvári, 2020, Bubeck and Cesa-Bianchi, 2012]. In this framework, a learning agent repeatedly chooses among a finite set of actions (called “arms”) and observes a noisy reward for the chosen arm, with the goal of maximizing cumulative reward over time. The appeal of the bandit model lies in its ability to capture the fundamental exploration-exploitation trade-off, that can be encountered in many sequential decision-making scenarios. Nevertheless, the classical bandit framework does not adequately capture two aspects that arise naturally in modern interactive learning systems: the dependence of rewards on user-specific contexts, and the potential cost of acquiring feedback.

An illustrative example is online content recommendation. Indeed, the quality of a recommendation depends crucially on the user who receives it: a video, news article, or product may be highly relevant to one user but uninteresting to another. This motivates the use of *contextual* bandit models [Abe and Long, 1999, Beygelzimer et al., 2011], where the expected reward depends on a context vector that describes the user or environment. A widely studied and practically successful instance is the linear contextual bandit model [Langford and Zhang, 2007, Li et al., 2010]. In this setting, the reward is modeled as the dot product between the observed context vector and an unknown arm-specific parameter. Linear contextual bandits offer a useful balance: they are expressive enough to capture heterogeneity in user preferences, while permitting efficient learning through regularized least-squares estimation.

A second challenge is that, in practice, feedback may not be observed automatically. While in standard bandits the learner always receives the reward of the chosen arm, in recommendation systems feedback often comes only if the user provides it (e.g., through ratings or explicit reviews). Actively requesting feedback at every round is undesirable, as it may burden or annoy users. A natural abstraction is therefore to associate a cost with each observation, so that the learner must strategically decide when feedback is worth acquiring. This leads to the framework of bandits with paid observations, first formalized by Seldin et al. [2014].

A third, orthogonal challenge is the nature of the reward-generating process. In some cases, user behavior is well modeled by a stochastic distribution, while in others it may be adversarial. Designing Best-of-Both-Worlds (BoBW) algorithms, that are versatile enough to perform optimally under both regimes, has become a central theme in bandit research [Bubeck and Slivkins, 2012, Zimmert and Seldin, 2022, Dann et al., 2023, Tsuchiya and Ito, 2024].

Motivated by these observations, in this work we introduce the setting of linear contextual bandits with paid observations, which simultaneously incorporates the challenges of contextual modeling, costly feedback acquisition, and uncertainty about the reward generation process. We design a new algorithm within the Follow-

the-Regularized-Leader (FTRL) framework, extending ideas from recent advances in best-of-both-worlds algorithms for bandits [Kuroki et al., 2024, Tsuchiya and Ito, 2024]. Our algorithm achieves regret guarantees in both stochastic and adversarial regimes, thereby solving the main challenges of the setting that we consider.

Achieving Best-of-Both-Worlds (BoBW) performance in hard problems, i.e. problems that incur a minimax regret of $\Theta(T^{2/3})$ in the adversarial regime, is a significant challenge, as highlighted in Tsuchiya and Ito [2024]. The standard approaches used in other settings often fail without substantial modifications. Fortunately, Tsuchiya and Ito [2024] introduced a dedicated framework designed to facilitate the design and analysis of BoBW algorithms for such problems. *#DB: the following will likely change after fixing the proof.* # While this framework forms the basis of our analysis, several challenges arise in adapting it to our setting. First, the general formulation assumes the existence of a single optimal arm throughout the learning process, which does not hold in the contextual linear setting where the optimal action varies with the context. Second, our setting introduces a new key parameter, the smallest non-negative eigenvalue of the context distribution (λ_{\min}), introduced in Section 2, which necessitates specific tuning of several algorithmic parameters. Third, we identify and resolve an inconsistency in prior applications of the BoBW framework to bandits with paid observations, thereby obtaining tighter regret guarantees; we elaborate on this point in Section 4. Structural differences in our setting require various other adjustments to the technical proofs.

1.1 Detailed literature review

In this section we detail existing results related to the different components of the settings that we consider.

Linear Contextual Bandits Contextual bandits extend classical multi-armed bandits by allowing the reward distribution to depend on an observed context, which can vary across rounds. To enable efficient decision-making, one must adopt a suitable model to capture how the context influences the rewards. In this work we consider the *linear contextual bandit* model [Langford and Zhang, 2007, Li et al., 2010], that we formally describe in Section 2. This model is closely-related to the well-studied *stochastic linear bandit* framework, since in both settings the average reward of each arm is given by the inner product of an arm feature vector and a parameter vector. The two formulations differ in the source of uncertainty: in stochastic linear bandits the arm features are known and the underlying parameter is unknown, whereas in (stochastic) linear contextual bandits the arm-specific

features are fixed but unknown, while the context vector is revealed at the beginning of each round.

Most approaches used in linear contextual bandits are borrowed from the stochastic linear bandit literature, in which algorithms follow general principles such as *Optimism in Face of Uncertainty* [Abe and Long, 1999, Dani et al., 2008, Abbasi-Yadkori et al., 2011, Flynn et al., 2023], *Thompson Sampling* [Agrawal and Goyal, 2013, Abeille and Lazaric, 2017, Abeille et al., 2025], *Information Directed Sampling* [Kirschner et al., 2020], or (asymptotic) lower bound matching [Lattimore and Szepesvári, 2017, Degenne et al., 2020]. Nonetheless, linear contextual bandits exhibit specific properties compared to standard linear bandits. In particular, Bastani et al. [2021] showed that under suitable assumptions on *context diversity*, even a simple greedy strategy can achieve logarithmic regret.

While the above works assume stochastic rewards, this assumption can be restrictive in practice. To address this, Neu and Olkhovskaya [2020] introduced an adversarial formulation of linear contextual bandits, in which arm parameters are fixed by an oblivious adversary. They derived a $\tilde{\mathcal{O}}(\sqrt{KdT})$ regret bound for an exponential-weights algorithm [Auer et al., 2002], where d is the parameter dimension, K is the number of arms, and T is the horizon. Building on this, Olkhovskaya et al. [2023] obtained refined first and second-order bounds. In parallel, Kuroki et al. [2024] established the first *Best-of-Both-Worlds* guarantees in this setting, showing that one can achieve simultaneously polylogarithmic regret in the stochastic regime and $\tilde{\mathcal{O}}(Kd\sqrt{T})$ regret in the adversarial case.

Bandits with Paid Observations This framework was introduced by Seldin and Slivkins [2014] to capture a feedback structure lying between the standard multi-armed bandit and full-information settings. In this model, the learner may choose to observe the reward of *any* arm at a fixed cost. They established that the minimax regret in this setting is $\Theta((cK)^{1/3}T^{2/3} + \sqrt{T})$, and proposed an algorithm matching this lower bound.

Prior to this, several related models were proposed to account for the possibility of observing additional feedback beyond the chosen arm [Mannor and Shamir, 2011, Avner et al., 2012, Alon et al., 2013], though these formulations do not explicitly capture the cost of information acquisition. An alternative approach is to impose a *budget* on the total observation cost, as in [Yun et al., 2018, Efroni et al., 2021]. However, this formulation requires the decision-maker to know both the acquisition cost of each arm and an overall budget, thereby placing regret minimization and acquisition costs on different scales. By contrast, the bandits-with-paid-observations framework integrates both aspects

under a unified metric by directly subtracting observation costs from the rewards.

Best-of-Both-Worlds (BoBW) The design of algorithms that perform well simultaneously in stochastic and adversarial regimes has become a central theme in the bandit literature. The foundational work of Bubeck and Slivkins [2012], Seldin and Slivkins [2014] initiated this line of research by asking whether one can achieve logarithmic regret in the stochastic setting while retaining $\tilde{O}(\sqrt{T})$ regret in the adversarial case. Their results provided only partial success, either with suboptimal bounds or with algorithms of limited practicality. Later, Zimmert and Seldin [2022] first obtained the optimal best-of-both-worlds guarantees in the K -armed bandit setting. This breakthrough has since inspired the development of BoBW algorithms across a variety of bandit problems [Amir et al., 2022, Rouyer et al., 2022, Saha and Gaillard, 2022, Tsuchiya et al., 2023, Jin et al., 2023, Zimmert and Marinov, 2024, Kato and Ito, 2025].

Of particular relevance to our work, Kuroki et al. [2024] studied linear contextual bandits through the black-box reduction framework of Dann et al. [2023], which can be used to design BoBW algorithms for problems whose minimax regret scales as \sqrt{T} . More recently, Tsuchiya and Ito [2024] proposed a general recipe for constructing BoBW algorithms in so-called “hard” online learning problems, namely those with minimax regret of order $\Theta(T^{2/3})$. They further show that several known bandit models, including multi-armed bandits with paid observations, fall within this framework. Our work is inspired by their approach, however, a direct application of their method does not yield optimal bounds in our setting (see Section 4). This motivates the need for a careful adaptation of their ideas, which we develop in the remainder of the paper.

2 PROBLEM DEFINITION

In this section we formalize the setting of *linear bandits with paid observations*, and state the main assumptions used in the analysis presented in Section 4.

Interaction protocol The interaction between the learning agent and the environment has a total duration of $T \in \mathbb{N}$ time steps, where T is unknown to the learner. Context vectors are drawn independently from a fixed distribution \mathcal{D} supported on a compact, full-dimensional subset $\mathcal{X} \subseteq \mathbb{R}^d$. At each round t , the following steps occur:

1. For each action $a \in [K] := 1, \dots, K$, the environment selects a loss parameter $\theta_{t,a} \in \mathbb{R}^d$.

2. A context $X_t \in \mathcal{X}$ is drawn from \mathcal{D} .
3. The learner observes X_t , chooses an action $A_t \in [K]$, and an observation set $O_t \subseteq [K]$.
4. The learner incurs loss $l_t(X_t, A_t) + c|O_t|$, where l_t is a loss function that depends on the environment parameters $(\theta_{t,a})_{a \in [K]}$, $c \in \mathbb{R}_{>0}$ is the known unit cost of observation, and $|O_t|$ is the cardinality of the observation set. It then observes the losses $\{l_t(X_t, o) : o \in O_t\}$.

Following Seldin and Slivkins [2014], the learner may query multiple arms in each round, paying cost c per queried arm. When $c = 0$, the learner is incentivized to query all arms, recovering the *full-information* (or “experts”) setting.

Assumptions To enable algorithm design and analysis, we adopt standard assumptions from the linear contextual bandit literature [Kuroki et al., 2024]:

1. $\|X\|_2 \leq 1$ almost surely.
2. $\forall t \in [T], a \in [K], \|\theta_{t,a}\|_2 \leq 1$.
3. $\forall t \in [T], x \in \mathcal{X}, a \in [K], l_t(x, a) \in [-1, 1]$.

We denote by $\Sigma = \mathbb{E}_{X \sim \mathcal{D}}[XX^\top] \succ 0$ the covariance matrix of the context distribution, and by $\lambda_{\min} > 0$ its minimum non zero eigenvalue, assumed to be known to the learner. While the learner does not know \mathcal{D} in full, we assume access to independent samples from \mathcal{D} between rounds, for instance through a simulator.

We now define how the loss $l_t(x, a)$ is constructed in each of the regimes considered in this work, for a given step $t \in [T]$, context $x \in \mathcal{X}$ and arm $a \in [K]$.

Adversarial regime The loss satisfies $l_t(X_t, a) := \langle X_t, \theta_{t,a} \rangle$, where $\theta_{t,a}$ is chosen by an *oblivious* adversary: the entire sequence $(\theta_{t,a})_{t \in [T], a \in [K]}$ can be arbitrary, but is fixed before the interaction starts.

Stochastic regime The loss is defined by $l_t(X_t, a) := \langle X_t, \theta_a \rangle + \varepsilon_{t,a}$ where θ_a is a fixed, unknown parameter for each arm a , and $\varepsilon_{t,a}$ is a zero-mean random noise bounded, independent across rounds and arms.

Corrupted stochastic regime The loss satisfies $l_t(X_t, a) := \langle X_t, \theta_{t,a} \rangle + \varepsilon_{t,a}$, where $\varepsilon_{t,a}$ is again a zero-mean random noise bounded in $[-1, 1]$. In this regime, the adversary may corrupt the parameters over time, but only within a limited budget: there exists fixed but unknown vectors $(\theta_a)_{a \in [K]}$ and a constant $C > 0$ such that $\sum_{t=1}^T \max_{a \in [K]} \|\theta_{t,a} - \theta_a\|_2 \leq C$. The

extreme cases $C = 0$ and $C = T$ recover, respectively, the stochastic regime and the adversarial regime (up to the presence of random noise).

Let Π denote the set of deterministic policies $\pi : \chi \mapsto [K]$. We define the best policy in hindsight π_T^* by

$$\pi_T^* : x \in \chi \longmapsto \arg \min_{a \in [K]} \mathbb{E} \left[\sum_{t=1}^T l_t(x, a) \right],$$

where potential randomness of the loss distribution. The learners' objective is to minimize the expected cumulative regret against π_T^* ,

$$R_T = \mathbb{E} \left[\sum_{t=1}^T (l_t(X_t, A_t) - l_t(X_t, \pi_T^*(X_t))) \right] + \mathbb{E} \left[\sum_{t=1}^T c \cdot |O_t| \right], \quad (1)$$

where the expectation here additionally includes the learner's internal randomization.

#DB: If we re-write the proof with ghost sample, explain here#

Additional definitions In the (corrupted) stochastic regime, we further define

$$\Delta_{\min}(x) := \min_{a \neq \pi_T^*(x)} \langle x, \theta_a - \theta_{\pi_T^*(x)} \rangle \quad \forall x \in \chi,$$

and the minimum sub-optimality gap

$$\Delta_{\min} := \min_{x \in \chi} \Delta_{\min}(x).$$

If the distribution \mathcal{D} over contexts is discrete, then Δ_{\min} is always strictly positive if all arms have distinct parameters. However, in the case where \mathcal{D} is continuous, it is possible that $\Delta_{\min} = 0$. In such cases, stochastic regret guarantees depending on Δ_{\min}^{-1} become vacuous. Nonetheless, the adversarial regret bounds remain valid regardless of the value of Δ_{\min} .

In the analysis, we denote by \mathcal{H}_t the filtration generated by all past contexts, actions, and observed losses. Finally, we use equivalently the notation $a = O(b)$ or $a \lesssim b$ when there exists a constant $\omega > 0$ such that $a \leq \omega b$, where ω is independent of the following problem-dependent quantities: $T, d, K, \Sigma, \mathcal{D}, C, \Delta_{\min}$.

3 ALGORITHM

As is standard in the best-of-both-worlds literature, our algorithm builds on the *Follow-the-Regularized-Leader* (FTRL) framework [see, e.g., Shalev-Shwartz, 2012, Sec. 2.3]. This general principle is characterized by three

key design choices: a *loss estimator*, a *learning-rate schedule*, and an appropriate *regularizer*.

To obtain loss estimates adapted to the linear contextual setting, we follow the approach of Kuroki et al. [2024], constructing importance-weighted regression estimates of the losses. For computational efficiency, we employ the *Matrix Geometric Resampling (MGR)* method [Neu and Bartók, 2013, Bartók et al., 2014, Kuroki et al., 2024], which guarantees tractability while controlling both the bias and variance of the estimates (see also Neu [2015]).

The other components of our algorithm are more directly inspired by Algorithm 2 of Tsuchiya and Ito [2024], which addresses the best-of-both-worlds problem for multi-armed bandits with paid observations. In particular, we adopt their use of a Tsallis entropy regularizer, an adaptive learning-rate schedule, and the computation of an *observation probability* that is uniform across arms. This probability is derived from the sampling probability vector produced by FTRL. This idea to use distinct observation and sampling probabilities originates from the initial work of [Seldin and Slivkins, 2014].

In the following we detail the components of our algorithm for linear contextual bandits with paid observations. The pseudo-code can be found in Algorithm 1.

Sampling distribution (FTRL) We recall that, at each round $t \geq 1$, the learner observes a context vector X_t , and must choose an action $A_t \in [K]$. As a first step, our algorithm computes a sampling distribution $q_t(\cdot | X_t) \in \Delta_K$, where Δ_K denotes the $K - 1$ -dimensional probability simplex. Following Tsuchiya and Ito [2024], this distribution is obtained through the *Follow-the-Regularized-Leader* (FTRL) principle, by solving the optimization problem

$$q_t(\cdot | X_t) \in \arg \min_{q \in \Delta_K} \left\{ \sum_{s=1}^{t-1} \langle q, \tilde{\ell}_s(X_t) \rangle + \psi_t(q) + \bar{\beta} H_{\alpha}(q) \right\} \quad (2)$$

This formulation involves the following components:

- **Loss estimates.** For each round $s \leq t - 1$,

$$\tilde{\ell}_s(X_t) := \left(\langle X_t, \tilde{\theta}_{s,1} \rangle, \dots, \langle X_t, \tilde{\theta}_{s,K} \rangle \right)^T, \quad (3)$$

where $\tilde{\theta}_{s,a}$ is an estimator of the linear loss parameter $\theta_{s,a} \in \mathbb{R}^d$ (see Eq. (5)).

- **Regularizer.** We use the Tsallis entropy, with

$$\psi_t(q) := -\frac{H_{\alpha}(q)}{\eta_t}, \quad \text{for } H_{\alpha}(q) := \frac{1}{\alpha} \sum_{a=1}^K (q_a^{\alpha} - q_a),$$

#DB: Should be $\frac{1}{\alpha-1}$, I see the error is propagated from [Tsuchiya and Ito \[2024\]](#) (below Eq. (10), but I guess in their derivations they then use the right one. # where $\eta_t > 0$ is the learning rate at time t , and we fix $\alpha := 1 - (\log K)^{-1}$. For convenience, we also define $\beta_t := 1/\eta_t$.

- **Additional parameters.** We set $\bar{\alpha} := 1 - \alpha$ and

$$\bar{\beta} := \frac{32Kd\sqrt{c}}{(1-\alpha)^2\sqrt{\beta_1}\min(1,\lambda_{\min})},$$

where c , K , and λ_{\min} are as introduced in Section 2. The term $\beta_1 = \eta_1^{-1}$ is introduced here in order to simplify some parts of the analysis, since we will define the learning rate such that $\beta_t \geq \beta_1$ holds for all time steps $t \geq 1$.

The definition of the FTRL distribution in Eq. (2) follows Algorithm 2 of [Tsuchiya and Ito \[2024\]](#), with two key modifications. The first, as previously discussed, is the use of loss estimates specifically adapted to the linear contextual structure of our setting.

The second is the value of $\bar{\beta}$ before the second regularization term, which we use in the analysis to control the evolution of $H_\alpha(q_t)$ between rounds (see Lemma 6), in particular at the beginning of the interaction (since this term doesn't scale up with t). This value is adjusted by the parameter λ_{\min} to account for the impact of the context distribution in the analysis.

Estimation of the linear losses We rely on a standard importance-weighted estimator, adapted from [Kuroki et al. \[2024\]](#). The key modification is that, instead of using the sampled action, we use the actions that are *observed* (if any) at round t . Specifically, for $t \geq 1$ and $a \in [K]$, we could estimate $\theta_{t,a}$ by

$$\hat{\theta}_{t,a} := \Sigma_{t,a}^{-1} X_t l_t(X_t, a) \mathbf{1}_{\{a \in O_t\}}, \quad (4)$$

where $\Sigma_{t,a} := \mathbb{E}[\mathbf{1}_{a \in O_t} X_t X_t^\top | \mathcal{H}_t]$. However, computing $\Sigma_{t,a}^{-1}$ exactly is computationally impractical for two reasons. First, matrix inversion at every round costs $\mathcal{O}(d^3)$ operations, which becomes prohibitive in high dimensions. Second, evaluating $\Sigma_{t,a}$ itself may be extremely costly: even in the discrete-context case, it requires computing observation probabilities for all possible contexts, with complexity at least $\mathcal{O}(|\mathcal{X}|)$, and moreover presupposes full knowledge of the context distribution.

To circumvent this issue, we approximate $\Sigma_{t,a}^{-1}$ using the *Matrix Geometric Resampling* (MGR) procedure, described in Algorithm 2 (Appendix). Computationally, MGR only requires sampling M_t contexts independently from \mathcal{D} , evaluating their observation probabilities (i.e., those the algorithm would assign if the context were observed at round t), and performing basic algebraic operations. This reduces the dependence

of the cost from $|\mathcal{X}|$ to $\mathcal{O}(\log(T))$, while only requesting access to a sampler of \mathcal{D} .

Accordingly, the estimator used in our algorithm is

$$\tilde{\theta}_{t,a} := \Sigma_{t,a}^+ X_t l_t(X_t, a) \mathbf{1}_{\{a \in O_t\}}, \quad (5)$$

where $\Sigma_{t,a}^+$ is the approximation of $\Sigma_{t,a}^{-1}$ returned by the MGR routine. Guided by our analysis, we set the number of MGR iterations to

$$M_t := \left\lceil \frac{4K}{p_t \lambda_{\min}} \ln(t) \right\rceil, \quad (6)$$

#DB: Will have to redefine this after fix# which ensures sufficiently accurate approximation of $\Sigma_{t,a}^+$. Compared to [Kuroki et al. \[2024\]](#), where the bias of the estimator is controlled via a forced exploration rate, in our setting this role is played by the observation probability p_t .

Observation probability Since observing each arm incurs a fixed cost c , the observation probability p_t must balance variance reduction with cost. We define

$$\begin{aligned} z_t &:= \frac{4cKd^2}{(1-\alpha)\lambda_{\min}^2} \left(q_{t*}^{2-\alpha} + \sum_{i \neq I_t} q_{ti}^{2-\alpha} \right), \\ u_t &:= \frac{8d \max(c, 1)}{(1-\alpha)\lambda_{\min}} q_{t*}^{1-\alpha}, \text{ where} \end{aligned} \quad (7)$$

$$I_t := \arg \max_{i \in [K]} q_{t,i}, \text{ and } q_{t*} := \min\{q_{t,I_t}, 1 - q_{t,I_t}\}.$$

#DB: d here introduced artificially, redo later.#

Compared to Algorithm 2 in [\[Tsuchiya and Ito, 2024\]](#), we have modified the definitions of the quantities z_t and u_t to include the λ_{\min} and d terms, which becomes necessary to appropriately control the variance of importance-weighted losses. For a learning rate η_t , we then define the observation probability as

$$p_t := \min \left\{ \frac{\sqrt{z_t \eta_t} + u_t \eta_t}{cK}, 1 \right\}. \quad (8)$$

This tuning seems to differ from the one proposed in Eq. 93 of [Tsuchiya and Ito \[2024\]](#) for their BoBW algorithm in the MAB with paid observations setting. As we explain in Section 4, our choice avoids a factor $(\frac{1}{cK} + cK)$ in the regret bound, which would otherwise render the guarantee vacuous when c is very small. Moreover, Eq. (7) shows that without this inverse scaling in c , the observation probability would converge to zero for small c under a fixed sampling probability, which is an unintuitive and undesirable behavior.

The fact that the probability p_t is uniform across arms has two important consequences for the MGR scheme. First, it removes the need for the forced exploration mechanism used in [\[Kuroki et al., 2024\]](#) to control the

bias (see their Lemma 9), and instead leads to a different result, formalized in our Lemma 8. Second, since $\Sigma_{t,a}$ is identical for all arms, we only need to compute a single pseudo-inverse Σ_t^+ per round. As a result, MGR only needs to be executed once at each time step, significantly reducing the overall computational cost.

Learning rate The learning rate η_t balances stability and adaptivity of FTRL, and is chosen to ensure optimal regret in both regimes. We follow Rule 2 of the framework presented in [Tsuchiya and Ito, 2024] and use the update rule

$$\frac{1}{\eta_{t+1}} = \frac{1}{\eta_t} + \frac{1}{h_t}(2\sqrt{z_t\eta_t} + u_t\eta_t), \quad (9)$$

where h_t denotes the entropy $H(q_t)$. For notational convenience we set $\gamma_t = cK \cdot p_t$. We also choose η_1 to ensure that $p_t \leq \frac{1}{2}$ for all time steps,

$$\eta_1 = \frac{(1-\alpha)\lambda_{\min}^2}{64\max(c,1)K} \quad (10)$$

Algorithm 1 FTRL for linear contextual bandits with paid observations

Require: K arms, cost c , λ_{\min} , $\forall a \in [K]$

- 1: Init η_1 as in (10), and $\forall a \in [K]$ set $\tilde{\theta}_{0,a} = 0$
- 2: **for** $t = 1, 2, \dots, T$ **do**
- 3: Observe X_t and compute $q_t(\cdot|X_t)$ as in (2)
- 4: sample $A_t \sim q_t(\cdot|X_t)$
- 5: Compute p_t as in (8)
- 6: For each $a \in [K]$, observe $l_t(X_t, a)$ with prob. p_t
- 7: Suffer the loss $l_t(X_t, A_t) + c|O_t|$
- 8: Update η_t to η_{t+1} according to (9)
- 9: $\forall a \in [K]$, compute and store $\tilde{\theta}_{t,a}$ via Alg. 2
- 10: Compute and store Σ_t^+ via MGR (see Algorithm 2) with M_t iterations.
- 11: **end for**

Computation time and memory The total space and time complexity of Algorithm 1 are respectively $\mathcal{O}(Td^2)$ and $\mathcal{O}(K^2T^2d^2\log T)$. Details can be found in Appendix F.

4 REGRET ANALYSIS

We now introduce the main theoretical result of this work, which is that Algorithm 1 achieves Best-of-Both-Worlds regret guarantees in the setting of linear bandits with paid observations, under the assumptions introduced in Section 2.

#DB: We have to change the dimension dependency after the fix. #

Theorem 1. In the adversarial regime, the regret of Algorithm 1 satisfies

$$R_T \lesssim \left(\frac{cKd^2 \log K}{\lambda_{\min}^2} \right)^{1/3} T^{2/3} + \sqrt{\frac{\max(c, 1)d \log K \cdot T}{\lambda_{\min}}} + \kappa$$

with

$$\kappa = \sqrt{\frac{cKd^2 \log K}{\lambda_{\min}^2}} + \frac{\max(c, 1)d \log K}{\lambda_{\min}} + \frac{\max(c, 1)K \log K}{\lambda_{\min}^2} + \frac{32Kd\sqrt{c}}{(1-\alpha)^2\sqrt{\beta_1}\min(1, \lambda_{\min})}.$$

while in the corrupted stochastic regime with corruption level C it satisfies

$$R_T \lesssim \frac{d\sqrt{\max(c, 1)K \log K}}{\lambda_{\min}\Delta_{\min}^2} \cdot \log(T\Delta_{\min}^3) + \left(\frac{C^2 d \sqrt{\max(c, 1)K \log K}}{\lambda_{\min}\Delta_{\min}^2} \cdot \log\left(\frac{T\Delta_{\min}}{C}\right) \right)^{1/3} + \kappa + \kappa', \text{ where we further define}$$

$$\kappa' = \left(\left(\frac{cKd^2 \log K}{\lambda_{\min}^2} \right)^{1/3} + \sqrt{\frac{\max(c, 1)d \log K}{\lambda_{\min}}} \right) \times \left(\frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}} \right)^{2/3}$$

This result shows that Algorithm 1 achieves the minimax-optimal $\mathcal{O}(T^{2/3})$ regret in the adversarial regime, while smoothly adapting to the (possibly corrupted) stochastic regime with logarithmic dependence on T when $C = 0$. These bounds match the known lower bounds from [Seldin et al., 2014], which applies to our setting since it encompasses the standard multi-armed bandit (by taking $d = 1$ and $X_t = 1$ a.s.), and extend the Best-of-Both-Worlds (BoBW) framework of [Tsuchiya and Ito, 2024] to the setting of linear bandits.

While the dependence in T is thus known to be optimal, the optimal dependence in other problem-specific parameters remains unknown, as this is the first work to address this setting. However, since our algorithm builds upon and generalizes both Algorithm 2 from [Kuroki et al., 2024] and Algorithm 2 from [Tsuchiya and Ito, 2024], we can compare our regret bounds to theirs, even if the settings do not perfectly align.

We consider first the limiting case where $c \rightarrow 0$, corresponding to the full-information setting, in which all losses are observed. In this regime, the first term of the adversarial regret bound vanishes, and we have

$$R_T \lesssim \sqrt{\frac{\log(K) \cdot dT}{\lambda_{\min}}}.$$

This matches, up to logarithmic factors, the adversarial regret bound established for Algorithm 2 in [Kuroki et al., 2024], namely

$$R_T \lesssim \sqrt{T \left(d + \frac{\log T}{\lambda_{\min}} \right) K \log K \log T}.$$

In our case, the factor K is replaced by $\log K$, which reflects the full-information nature of our setting, a standard improvement in such regimes. However, in the stochastic regime, our regret exhibits an additional $\frac{1}{\Delta_{\min}}$ factor compared to the full-information bounds in [Kuroki et al., 2024]. But on the contrary, our algorithm has a better $\log T$ dependence, thus our bound is better if T is significantly larger than $\frac{1}{\Delta_{\min}}$. Although, we do not know whether our improved $\log T$ dependency stems from being in the full-information setting or from other factors. We can at least observe that the dependence on the setting-specific parameters d and λ_{\min} in our bounds matches that of their Algorithm 2.

Another useful comparison is to consider the special case $d = 1, \mathcal{X} = \{1\}$, in which case we recover the setting of Seldin and Slivkins [2014]. From their Corollary 17, Algorithm 2 of Tsuchiya and Ito [2024] obtain an adversarial regret bound of

$$\mathcal{R}_T \lesssim ((cK)^{1/3} T^{2/3} (\log K)^{1/3}),$$

which is exactly the scaling that we obtain with Theorem 1 in this setting. This observation furthermore still holds in the stochastic setting.

These comparisons suggest that, while we can not establish optimality in general due to the lack of known lower bounds, our algorithm can be viewed as a strict generalization of the approach in [Tsuchiya and Ito, 2024] for bandits with paid observations, since we recover their guarantees in this setting. Moreover, since the dependencies in d and λ_{\min} are known to be optimal compared to previous approaches when $c = 0$, this further supports the relevance of our design beyond prior approaches.

A detailed proof of the theorem can be found in Appendix C. In the following, we present the main steps of the proofs, highlighting the technical arguments that required to be adapted from the existing frameworks.

Proof sketch. As a preliminary step of the analysis, we isolate the difficulty induced by the use of (biased) MGR estimates (Eq. (5)) instead of using the unbiased estimators from Eq. (3). Following the proof technique of Kuroki et al. [2024], we introduce an auxiliary game where these estimators are treated as unbiased, and for

which the regret would thus become

$$\tilde{R}_T := \mathbb{E} \left[\sum_{t=1}^T \left\langle X_t, \tilde{\theta}_{t,A_t} \right\rangle - \left\langle X_t, \tilde{\theta}_{t,\pi^*(X_t)} \right\rangle \right].$$

We can verify that the actual regret of our algorithm thus satisfies

$$R_T \leq \tilde{R}_T + 2 \sum_{t=1}^T \max_{a \in [K]} \left| \mathbb{E} \left[\left\langle X_t, \tilde{\theta}_{t,a} - \theta_{t,a} \right\rangle \right] \right|.$$

Then, in Lemma 9 we prove that the second term of this upper bound can be upper bounded by a constant, independent of all problem parameters. In the following, we thus focus on upper bounding \tilde{R}_T . We write the following proof steps with the notation R_T , with an abuse of notation, since previous result showed that both terms have the same scaling in T .

The remainder of the analysis builds on the general framework introduced by Tsuchiya and Ito [Tsuchiya and Ito, 2024] to build Best-of-Both-Worlds algorithms for problems with minimax regret scaling with $T^{2/3}$, and in particular their instantiation of this framework to tackle standard multi-armed bandit with paid observations (without the linear contextual structure). Our first contribution is an adaptation of their Theorem 7 to accommodate the linear contextual structure, that we introduce below.

Lemma 1 (Adaptation of Theorem 7 of [Tsuchiya and Ito, 2024]). *Suppose that Algorithm 1 satisfies the following conditions in the adversarial regime:*

- (i) $R_T \leq \sum_{t=1}^T \mathbb{E} \left[\left(\frac{1}{\eta_t} - \frac{1}{\eta_{t-1}} \right) h_t + \frac{z_t \eta_t}{\gamma_t} + \gamma_t \right] + \bar{\beta} h,$
- (ii) $\mathbb{E}[h_{t+1} | \mathcal{H}_t] \leq 2 \mathbb{E}[h_t | \mathcal{H}_{t-1}]$ for all $t \geq 1$.

Then the regret can be bounded as

$$R_T \lesssim (z_{\max} h_1)^{1/3} T^{2/3} + \sqrt{u_{\max} h_1 T} + \kappa,$$

where

$$z_{\max} = \max_{t \in [T]} z_t \leq 4cK \log K \frac{1}{\lambda_{\min}^2},$$

$$u_{\max} = \max_{t \in [T]} u_t \leq 4 \max(c, 1) \log K \frac{1}{\lambda_{\min}},$$

and

$$\kappa := \sqrt{z_{\max} \eta_1} + u_{\max} \eta_1 + \frac{h_1}{\eta_1} + \bar{\beta} h_{\max}.$$

Moreover, if Algorithm 1 satisfies the following conditions in the stochastic regime: there exists a constant $\rho > 0$ such that, $\forall t \geq 1$,

(iii) $\sqrt{z_t h_t} \leq \sqrt{\rho} (1 - \pi_T^*(X_t) | X_t)$, and

(iv) $u_t h_t \leq \rho (1 - \pi_T^*(X_t) | X_t)$,

then, for $T \geq \tau := \frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}}$ it holds that

$$R_T \lesssim \frac{\rho}{\Delta_{\min}^2} \log(T\Delta_{\min}^3) + \left(\frac{C^2 \rho}{\Delta_{\min}^2} \log(\frac{T\Delta_{\min}}{C}) \right)^{1/3} + \kappa'$$

with

$$\kappa' := \kappa + ((z_{\max} h_1)^{1/3} + \sqrt{u_{\max} h_1}) \left(\frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}} \right)^{2/3}.$$

While Lemma 1 adapts Theorem 7 from [Tsuchiya and Ito, 2024], it differs in several significant aspects. First, condition (i) is new and replaces conditions (i)–(ii) in the original theorem, and both lead to a similar proof structure, our condition better adjust the framework to our setting. Second, condition (ii) is a relaxed reformulation of condition (iii) in [Tsuchiya and Ito, 2024], which is necessary to handle the stochasticity of contexts in our setting. With careful use of the tower rule, we show that this weaker assumption is sufficient for the regret analysis. Finally, conditions (iii) and (iv) are reformulations of conditions (iv) and (v) from [Tsuchiya and Ito, 2024], and the corresponding proof techniques carry over with only little modifications. The detailed proof of this lemma is deferred to Appendix B.

To establish Theorem 1, it then suffices to verify that Algorithm 1 satisfies each of the four conditions.

Condition (i) follows from the standard FTRL regret decomposition: the stability term bound is direct to obtain, while the penalty term is controlled using Lemma 3 (in Appendix), which is similarly to the proof of [Tsuchiya and Ito, 2024, Theorem 8].

We prove condition (ii) in Lemma 6. The proof consists in applying Lemma 15 from [Tsuchiya and Ito, 2024] (restated as Lemma 5) for each fixed context, and to conclude via linearity of expectation. A key challenge arises from the fact that, in our setting, we have the bound $\mathbb{E} [\langle X_t, \hat{\theta}_{t,a} \rangle]^2 \leq \frac{1}{\lambda_{\min}^2 p_t}$, which contrasts with the original bound $\mathbb{E} [l_t^2] \leq \frac{1}{p_t}$ in the non-contextual case. Since Lemma 5 only accommodates a constant upper bound, this discrepancy required a careful adjustment of several parameters, specifically u_t and $\bar{\beta}$, which represents a slight modification in the precise behavior of the algorithm.

Finally, Conditions (iii) and (iv) are verified by combining entropy bounds from [Tsuchiya and Ito, 2024] with direct control of the variance-like quantities z_t and u_t , thereby linking them to the optimal action probability.

Together, these arguments ensure that Algorithm 1 satisfies the assumptions of Lemma 1, which directly yields the regret guarantees stated in Theorem 1.

The full derivations and supporting lemmas are deferred to Appendix C, where we carefully establish that each condition of the lemma holds in our setting. \square

While the definition of p_t in [Tsuchiya and Ito, 2024] differs from ours by a factor $(cK)^{-1}$, this appears to be a simple typo in their presentation. Indeed, their analysis assumes $p_t = \frac{1}{cK} (\sqrt{z_t \eta_t} + u_t \eta_t)$, even though the statement of their Algorithm 2 defines $p_t := \sqrt{z_t \eta_t} + u_t \eta_t$. We can use this observation to comment on the optimality of the tuning of p_t with respect to the analysis used to derive BoBW regret bounds for our algorithm. Indeed, a step in the analysis (see Eq. (11)) involves the quantity $\gamma'_t := \gamma_t - \frac{u_t}{\beta_t}$. With our definition, this yields $\gamma'_t = \sqrt{z_t / \beta_t}$, while using the unnormalized p_t (without $1/(cK)$) gives

$$\gamma'_t = cK \sqrt{z_t \eta_t} + (cK - 1) u_t \eta_t \geq cK \sqrt{z_t \eta_t},$$

assuming $cK \geq 1$. This leads to the bound

$$\begin{aligned} \sum_{t=1}^T \mathbb{E} \left[\frac{z_t \eta_t}{\gamma'_t} + \gamma_t \right] &\leq \sum_{t=1}^T \mathbb{E} \left[\frac{1}{cK} \sqrt{\frac{z_t}{\beta_t}} + cK \left(\sqrt{\frac{z_t}{\beta_t}} + \frac{u_t}{\beta_t} \right) \right] \\ &\leq \left(\frac{1}{cK} + cK \right) \sum_{t=1}^T \mathbb{E} \left[2 \sqrt{\frac{z_t}{\beta_t}} + \frac{u_t}{\beta_t} \right]. \end{aligned}$$

The factor $(cK)^{-1} + cK$ then propagates through the analysis and degrades the regret bound. More generally, an overestimation of p_t by a multiplicative factor ω leads to a regret that is worsened by a factor proportional to $\omega + \omega^{-1}$, so $\omega = 1$ (our tuning) is optimal.

5 DISCUSSION

We proposed an algorithm achieving BoBW regret guarantees in the setting of *linear contextual bandits with paid observations*, with explicit scaling in problem dimensions (d, K) and parameters ($\lambda_{\min}, \Delta_{\min}, c$).

However, an important limitation, shared with the analysis of Algorithm 2 from [Kuroki et al., 2024], arises in the stochastic setting when the context space is continuous. In such cases, the quantity Δ_{\min} is often zero, which implies that the regret bound remains at $\Theta(T^{2/3})$, even though the environment is stochastic and should, in principle, allow for better rates. This issue also affects discrete but finely spaced context spaces, where $\Delta_{\min} > 0$ but can be arbitrarily small, leading to overly pessimistic bounds in practice. Nevertheless, [Bastani et al., 2021] demonstrates that under suitable regularity conditions on the context distribution, it is possible to achieve logarithmic regret in continuous settings without any dependence on Δ_{\min} . Extending such ideas to our setting, and combining them with BoBW-style guarantees, could lead to improved regret

bounds, potentially polylogarithmic or polynomially better than \sqrt{T} or $T^{2/3}$. We believe this is a promising direction for future work.

Finally, as previously discussed, since this setting is novel, there are currently no lower bounds specifically tailored to it. Existing lower bounds only apply to simplified or special cases of our setting. Developing minimax and stochastic lower bounds that are adapted to this setting, precisely capturing all dimensions and parameters, would therefore be an interesting contribution to improve the understanding of this setting.

References

- Y. Abbasi-Yadkori, D. Pál, and C. Szepesvári. Improved algorithms for linear stochastic bandits. In *Advances in Neural Information Processing Systems 24*, pages 2312–2320, 2011.
- N. Abe and P. M. Long. Associative reinforcement learning using linear probabilistic concepts. In *Proceedings of the Sixteenth International Conference on Machine Learning*, pages 3–11. Morgan Kaufmann, 1999.
- M. Abeille and A. Lazaric. Linear thompson sampling revisited. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54, pages 176–184. PMLR, 2017.
- M. Abeille, D. Janz, and C. Pike-Burke. When and why randomised exploration works (in linear bandits). In *Proceedings of the International Conference on Algorithmic Learning Theory*, volume 272, pages 4–22. PMLR, 2025.
- S. Agrawal and N. Goyal. Thompson sampling for contextual bandits with linear payoffs. In *Proceedings of the 30th International Conference on Machine Learning*, pages 127–135. JMLR.org, 2013.
- N. Alon, N. Cesa-Bianchi, C. Gentile, and Y. Mansour. From bandits to experts: A tale of domination and independence. In *Advances in Neural Information Processing Systems 26*, pages 1610–1618, 2013.
- I. Amir, G. Azov, T. Koren, and R. Livni. Better best of both worlds bounds for bandits with switching costs. *Advances in Neural Information Processing Systems*, 35:15800–15810, 2022.
- P. Auer, N. Cesa-Bianchi, Y. Freund, and R. E. Schapire. The nonstochastic multiarmed bandit problem. *SIAM Journal on Computing*, 32(1):48–77, 2002.
- O. Avner, S. Mannor, and O. Shamir. Decoupling exploration and exploitation in multi-armed bandits. In *Proceedings of the 29th International Conference on Machine Learning*, 2012.
- G. Bartók, D. Foster, D. Pál, A. Rakhlin, and C. Szepesvári. Partial monitoring—classification, regret bounds, and algorithms. *Mathematics of Operations Research*, 39(4):967–997, 2014. doi: 10.1287/moor.2014.0663.
- H. Bastani, M. Bayati, and K. Khosravi. Mostly exploration-free algorithms for contextual bandits. *Management Science*, 67(3):1329–1349, 2021. doi: 10.1287/mnsc.2020.3605.
- A. Beygelzimer, J. Langford, L. Li, L. Reyzin, and R. E. Schapire. Contextual bandit algorithms with supervised learning guarantees. In *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, volume 15, pages 19–26. JMLR.org, 2011.
- S. Bubeck and N. Cesa-Bianchi. Regret analysis of stochastic and nonstochastic multi-armed bandit problems. *Foundations and Trends in Machine Learning*, 5(1):1–122, 2012. doi: 10.1561/2200000024.
- S. Bubeck and A. Slivkins. The best of both worlds: Stochastic and adversarial bandits. In *Proceedings of the 25th Annual Conference on Learning Theory*, volume 23, pages 42.1–42.23. PMLR, 2012.
- V. Dani, T. P. Hayes, and S. M. Kakade. Stochastic linear optimization under bandit feedback. In *Proceedings of the 21st Annual Conference on Learning Theory*, pages 355–366. Omnipress, 2008.
- C. Dann, C.-Y. Wei, and J. Zimmert. A blackbox approach to best of both worlds in bandits and beyond. In *Proceedings of the 36th Annual Conference on Learning Theory*, volume 195, pages 5503–5570. PMLR, 2023.
- R. Degenne, H. Shao, and W. M. Koolen. Structure adaptive algorithms for stochastic bandits. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119, pages 2443–2452. PMLR, 2020.
- Y. Efroni, N. Merlis, A. Saha, and S. Mannor. Confidence-budget matching for sequential budgeted learning. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139, pages 2937–2947. PMLR, 2021.
- H. Flynn, D. Reeb, M. Kandemir, and J. R. Peters. Improved algorithms for stochastic linear bandits using tail bounds for martingale mixtures. In *Advances in Neural Information Processing Systems 36*, 2023.
- T. Jin, J. Liu, and H. Luo. Improved best-of-both-worlds guarantees for multi-armed bandits: Ftrl with general regularizers and multiple optimal arms. *Advances in Neural Information Processing Systems*, 36: 30918–30978, 2023.

- M. Kato and S. Ito. Lc-tsallis-inf: Generalized best-of-both-worlds linear contextual bandits. In *Proceedings of the International Conference on Artificial Intelligence and Statistics*, volume 258, pages 3655–3663. PMLR, 2025.
- J. Kirschner, T. Lattimore, and A. Krause. Information directed sampling for linear partial monitoring. In *Proceedings of the Conference on Learning Theory*, volume 125, pages 2328–2369. PMLR, 2020.
- Y. Kuroki, A. Rumi, T. Tsuchiya, F. Vitale, and N. Cesa-Bianchi. Best-of-both-worlds algorithms for linear contextual bandits. In S. Dasgupta, S. Mandt, and Y. Li, editors, *Proceedings of the 27th International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings of Machine Learning Research*, pages 1216–1224. PMLR, 2024.
- J. Langford and T. Zhang. The epoch-greedy algorithm for multi-armed bandits with side information. In *Advances in Neural Information Processing Systems 20*, pages 817–824, 2007.
- T. Lattimore and C. Szepesvári. The end of optimism? an asymptotic analysis of finite-armed linear bandits. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54, pages 728–737. PMLR, 2017.
- T. Lattimore and C. Szepesvári. *Bandit Algorithms*. Cambridge University Press, 2020.
- L. Li, W. Chu, J. Langford, and R. E. Schapire. A contextual-bandit approach to personalized news article recommendation. In *Proceedings of the 19th International Conference on World Wide Web*, pages 661–670. ACM, 2010. doi: 10.1145/1772690.1772758.
- S. Mannor and O. Shamir. From bandits to experts: On the value of side-observations. In *Advances in Neural Information Processing Systems 24*, pages 684–692, 2011.
- G. Neu. Explore no more: Improved high-probability regret bounds for non-stochastic bandits. *arXiv preprint arXiv:1506.03271*, 2015.
- G. Neu and G. Bartók. An efficient algorithm for learning with semi-bandit feedback. *arXiv preprint arXiv:1305.2732*, 2013.
- G. Neu and J. Olkhovskaya. Efficient and robust algorithms for adversarial linear contextual bandits. In *Proceedings of the Conference on Learning Theory*, volume 125, pages 3049–3068. PMLR, 2020.
- J. Olkhovskaya, J. J. Mayo, T. van Erven, G. Neu, and C.-Y. Wei. First- and second-order bounds for adversarial linear contextual bandits. In *Advances in Neural Information Processing Systems 36*, 2023.
- C. Rouyer, D. van der Hoeven, N. Cesa-Bianchi, and Y. Seldin. A near-optimal best-of-both-worlds algorithm for online learning with feedback graphs. In *Advances in Neural Information Processing Systems 35*, 2022.
- A. Saha and P. Gaillard. Versatile dueling bandits: Best-of-both world analyses for learning from relative preferences. In *Proceedings of the International Conference on Machine Learning*, pages 19011–19026. PMLR, 2022.
- Y. Seldin and A. Slivkins. One practical algorithm for both stochastic and adversarial bandits. In *Proceedings of the 31st International Conference on Machine Learning*, pages 1287–1295. JMLR.org, 2014.
- Y. Seldin, P. Bartlett, K. Crammer, and Y. Abbasi-Yadkori. Prediction with limited advice and multi-armed bandits with paid observations. In E. Xing and T. Jebara, editors, *Proceedings of the 31st International Conference on Machine Learning*, volume 32 of *Proceedings of Machine Learning Research*, pages 280–287, Beijing, China, 2014. PMLR.
- Shai Shalev-Shwartz. Online learning and online convex optimization. *Foundations and Trends in Machine Learning*, 4(2):107–194, 2012.
- T. Tsuchiya and S. Ito. A simple and adaptive learning rate for ftrl in online learning with minimax regret of $\theta(t^{2/3})$ and its application to best-of-both-worlds. *NeurIPS*, 2024.
- T. Tsuchiya, S. Ito, and J. Honda. Further adaptive best-of-both-worlds algorithm for combinatorial semi-bandits. In *Proceedings of the International Conference on Artificial Intelligence and Statistics*, pages 8117–8144. PMLR, 2023.
- D. Yun, A. Proutière, S. Ahn, J. Shin, and Y. Yi. Multi-armed bandit with additional observations. *Proceedings of ACM Measurement and Analysis of Computing Systems*, 2(1):13:1–13:22, 2018. doi: 10.1145/3179416.
- J. Zimmert and T. V. Marinov. Productive bandits: Importance weighting no more. In *Advances in Neural Information Processing Systems 37*, pages 85360–85388, 2024.
- J. Zimmert and Y. Seldin. Tsallis-inf: An optimal algorithm for stochastic and adversarial bandits, 2022.

Checklist

1. For all models and algorithms presented, check if you include:
 - (a) A clear description of the mathematical setting, assumptions, algorithm, and/or model.

- [Yes] The setting and assumptions are presented in Section 2, while the algorithm is presented in Section 3.
- (b) An analysis of the properties and complexity (time, space, sample size) of any algorithm. [Yes] A time and space complexity analysis is provided at the end of Section 3, while Section 4 offers a detailed analysis of the regret.
 - (c) (Optional) Anonymized source code, with specification of all dependencies, including external libraries. [Not Applicable]
2. For any theoretical claim, check if you include:
- (a) Statements of the full set of assumptions of all theoretical results. [Yes] The main theoretical result is stated as Theorem 1.
 - (b) Complete proofs of all theoretical results. [Yes] The proof of Theorem 1 is explained in Section 4 and detailed throughout the Appendix.
 - (c) Clear explanations of any assumptions. [Yes] All assumptions are explained in detail in Section 2.
3. For all figures and tables that present empirical results, check if you include:
- (a) The code, data, and instructions needed to reproduce the main experimental results (either in the supplemental material or as a URL). [Not Applicable]
 - (b) All the training details (e.g., data splits, hyperparameters, how they were chosen). [Not Applicable]
 - (c) A clear definition of the specific measure or statistics and error bars (e.g., with respect to the random seed after running experiments multiple times). [Not Applicable]
 - (d) A description of the computing infrastructure used. (e.g., type of GPUs, internal cluster, or cloud provider). [Not Applicable]
4. If you are using existing assets (e.g., code, data, models) or curating/releasing new assets, check if you include:
- (a) Citations of the creator If your work uses existing assets. [Not Applicable]
 - (b) The license information of the assets, if applicable. [Not Applicable]
 - (c) New assets either in the supplemental material or as a URL, if applicable. [Not Applicable]
 - (d) Information about consent from data providers/curators. [Not Applicable]
- (e) Discussion of sensible content if applicable, e.g., personally identifiable information or offensive content. [Not Applicable]
5. If you used crowdsourcing or conducted research with human subjects, check if you include:
- (a) The full text of instructions given to participants and screenshots. [Not Applicable]
 - (b) Descriptions of potential participant risks, with links to Institutional Review Board (IRB) approvals if applicable. [Not Applicable]
 - (c) The estimated hourly wage paid to participants and the total amount spent on participant compensation. [Not Applicable]

A To-Do

#DB: This section is for me to list the mistakes I remarked in the proof and some thoughts from related work. #

- #DB: We have to cleanly re-write the proof with the ghost sample technique#, which is probably going to simplify some things (e.g. Lemma 6, maybe the property is only needed for the ghost context). To do that, work directly from Lemma 3 of Neu and Olkhovskaya [2020]. We will probably be able to get the dimension from the standard trace trick.
- In Eq. (12) if we want to use Lemma 2 we have to upper bounds the probs by 1 and get K instead of $\sum q_a^{2-\alpha}$. Otherwise we need a bound on $\mathbb{E}[q_{t,a}^{2-\alpha} \hat{\ell}_{t,a}^2]$ directly.
- In the MGR concentration and the proof of Lemma 2 we wrongly do as if the observation probability was uniform across all arms, the only way to make it right seems to use $\min_{X \in \mathcal{X}} p_t(X)$. #DB: Several errors of this type can be corrected by introducing a forced exploration.#

In addition, below are some things that I don't understand and some remarks

- In those papers they also use what is our Lemma 2 on the MGR-induced losses directly, not on the "unbiased estimates losses".
- I really want to understand what "should" be the right dependency in d , if not from a formal lower bound at least from an intuitive point of view.
- Kuroki et al. have a worst stochastic bound in $(\log(T))^3 + d(\log(T))^2$. I'm wondering why and if we should have the same or if Taira's framework prevents us from that.

Ghost samples Adapting the proofs from previous works we can easily get that

$$R_T \leq \mathbb{E}_{X_0 \sim \mathcal{D}} [\tilde{R}_T(X_0)] + 2 \sum_{t=1}^T \text{bias-MGR}_t,$$

where for any $x \in \mathcal{D}$ we define

$$\tilde{R}_T(x) = \sum_{t=1}^T \left\{ \sum_{a=1}^K (q_t(a|x) - \pi_T^\star(a|x)) \langle x, \tilde{\theta}_{t,a} \rangle + cKp_t(x) \right\},$$

where $p_t(x)$ is the observation probability at time t if context x is revealed.

#DB: To me it is on $\tilde{R}_t(x)$ that we have to state and prove Lemma 1.#

B PROOF OF LEMMA 1

Lemma 1 (Adaptation of Theorem 7 of [Tsuchiya and Ito, 2024]). *Suppose that Algorithm 1 satisfies the following conditions in the adversarial regime:*

$$(i) R_T \leq \sum_{t=1}^T \mathbb{E} \left[\left(\frac{1}{\eta_t} - \frac{1}{\eta_{t-1}} \right) h_t + \frac{z_t \eta_t}{\gamma_t} + \gamma_t \right] + \bar{\beta} \bar{h},$$

$$(ii) \mathbb{E}[h_{t+1} | \mathcal{H}_t] \leq 2 \mathbb{E}[h_t | \mathcal{H}_{t-1}] \text{ for all } t \geq 1.$$

Then the regret can be bounded as

$$R_T \lesssim (z_{\max} h_1)^{1/3} T^{2/3} + \sqrt{u_{\max} h_1 T} + \kappa,$$

where

$$\begin{aligned} z_{\max} &= \max_{t \in [T]} z_t \leq 4cK \log K \frac{1}{\lambda_{\min}^2}, \\ u_{\max} &= \max_{t \in [T]} u_t \leq 4 \max(c, 1) \log K \frac{1}{\lambda_{\min}}, \end{aligned}$$

and

$$\kappa := \sqrt{z_{\max} \eta_1} + u_{\max} \eta_1 + \frac{h_1}{\eta_1} + \bar{\beta} h_{\max}.$$

Moreover, if Algorithm 1 satisfies the following conditions in the stochastic regime: there exists a constant $\rho > 0$ such that, $\forall t \geq 1$,

(iii) $\sqrt{z_t h_t} \leq \sqrt{\rho} (1 - \pi_T^*(X_t) | X_t)$, and

(iv) $u_t h_t \leq \rho (1 - \pi_T^*(X_t) | X_t)$,

then, for $T \geq \tau := \frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}}$ it holds that

$$R_T \lesssim \frac{\rho}{\Delta_{\min}^2} \log(T \Delta_{\min}^3) + \left(\frac{C^2 \rho}{\Delta_{\min}^2} \log(\frac{T \Delta_{\min}}{C}) \right)^{1/3} + \kappa'$$

with

$$\kappa' := \kappa + ((z_{\max} h_1)^{1/3} + \sqrt{u_{\max} h_1}) \left(\frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}} \right)^{2/3}.$$

Proof of Lemma 1. The argument follows the same general structure of the proof of Theorem 7 in [Tsuchiya and Ito \[2024\]](#). We first define

$$\gamma'_t := \gamma_t - \frac{u_t}{\beta_t} = \sqrt{\frac{z_t}{\beta_t}}.$$

Starting from the regret decomposition given by Assumption (i) of the lemma, we have:

$$\begin{aligned} R_T &\leq \mathbb{E} \left[\sum_{t=1}^T \left(\left(\frac{1}{\eta_t} - \frac{1}{\eta_{t-1}} \right) h_t + \frac{z_t \eta_t}{\gamma_t} + \gamma_t \right) \right] + \bar{\beta} \bar{h} \\ &\quad (\text{by Assumption (i) of the lemma}) \\ &= \mathbb{E} \left[\sum_{t=1}^T \left((\beta_t - \beta_{t-1}) h_t + \frac{z_t \eta_t}{\gamma_t} + \gamma_t \right) \right] + \bar{\beta} \bar{h} \end{aligned}$$

where we used $\beta_t = 1/\eta_t$ and $\bar{h} = \max_{p \in \Delta_K} H_{\bar{\alpha}}(p)$. We now replace γ_t by $\gamma'_t \leq \gamma_t$ to simplify the analysis (this may loosen the bound slightly but keeps the algebra tractable):

$$\begin{aligned} R_T &\leq \mathbb{E} \left[\sum_{t=1}^T \left((\beta_t - \beta_{t-1}) h_t + \frac{z_t \eta_t}{\gamma'_t} + \gamma_t \right) \right] + \bar{\beta} \bar{h} \\ &= \sum_{t=1}^T \mathbb{E} \left[(\beta_t - \beta_{t-1}) h_t + \frac{z_t \eta_t}{\gamma'_t} + \gamma_t \right] + \bar{\beta} \bar{h} \end{aligned}$$

Bounding the first term. We first upper bound $\sum_{t=1}^T \mathbb{E}[(\beta_t - \beta_{t-1}) h_t]$. By the tower rule and the fact that $(\beta_t - \beta_{t-1})$ is \mathcal{H}_{t-1} -measurable, we have

$$\sum_{t=1}^T \mathbb{E}[(\beta_t - \beta_{t-1}) h_t] = \sum_{t=1}^T \mathbb{E}[(\beta_t - \beta_{t-1}) \mathbb{E}[h_t | \mathcal{H}_{t-1}]].$$

Then, by Assumption (ii), which ensures $\mathbb{E}[h_t | \mathcal{H}_{t-1}] \leq 2 \mathbb{E}[h_{t-1} | \mathcal{H}_{t-2}]$, we get:

$$\sum_{t=1}^T \mathbb{E}[(\beta_t - \beta_{t-1}) h_t] \leq 2 \sum_{t=1}^T \mathbb{E}[(\beta_t - \beta_{t-1}) h_{t-1}].$$

Bounding the remaining terms. Using the definitions of γ_t and γ'_t , namely $\gamma'_t = \sqrt{z_t/\beta_t}$ and $\gamma_t = \gamma'_t + u_t/\beta_t$, we have

$$\begin{aligned} \sum_{t=1}^T \mathbb{E} \left[\frac{z_t \eta_t}{\gamma'_t} + \gamma_t \right] &= \sum_{t=1}^T \mathbb{E} \left[\sqrt{\frac{z_t}{\beta_t}} + \left(\sqrt{\frac{z_t}{\beta_t}} + \frac{u_t}{\beta_t} \right) \right] \\ &\leq \sum_{t=1}^T \mathbb{E} \left[2 \sqrt{\frac{z_t}{\beta_t}} + \frac{u_t}{\beta_t} \right]. \end{aligned} \quad (11)$$

Combining the two results above, we obtain:

$$R_T \leq \mathbb{E}[F(\beta_{1:T}, z_{1:T}, u_{1:T}, h_{0:T-1})] + \mathbb{E}[\bar{\beta}h],$$

where we define

$$F(\beta_{1:T}, z_{1:T}, u_{1:T}, h_{1:T}) := \sum_{t=1}^T \left((\beta_t - \beta_{t-1})h_t + 2\sqrt{\frac{z_t}{\beta_t}} + \frac{u_t}{\beta_t} \right).$$

Note that the regret upper bound we obtained at this step involves the sequence $h_{0:T-1}$, and not $h_{1:T}$ as in the above definition.

Adversarial regime Using Lemma 7, we obtain that, for any $\varepsilon \geq 1/T$,

$$\begin{aligned} F(\beta_{1:T}, z_{1:T}, u_{1:T}, h_{0:T-1}) &\leq \left(\left[\sum_{t=1}^T \sqrt{z_t h_t} \right] \log(\varepsilon T) \right)^{2/3} \\ &\quad + \sqrt{\left[\sum_{t=1}^T u_t h_t \right] \log(\varepsilon T)} \\ &\quad + \left(\frac{\sqrt{z_{\max} h_1}}{\varepsilon} \right)^{2/3} + \sqrt{\frac{u_{\max} h_1}{\varepsilon}} + \kappa. \end{aligned}$$

Substituting this into the previous inequality gives the claimed regret bound for the adversarial case:

$$\begin{aligned} R_T &\leq \left(\mathbb{E} \left[\sum_{t=1}^T \sqrt{z_t h_t} \right] \log(\varepsilon T) \right)^{2/3} + \sqrt{\mathbb{E} \left[\sum_{t=1}^T u_t h_t \right] \log(\varepsilon T)} \\ &\quad + \left(\frac{\sqrt{z_{\max} h_1}}{\varepsilon} \right)^{2/3} + \sqrt{\frac{u_{\max} h_1}{\varepsilon}} + \kappa. \end{aligned}$$

Setting $\varepsilon = 1/T$ yields the desired bound in the adversarial regime.

Stochastic regime. We now turn to the stochastic case, under Assumptions (iii)–(iv). Define

$$\varrho_0(\pi_T^\star) := \sum_{t=1}^T (1 - q_t(\pi_T^\star(X_t) | X_t)).$$

By Assumptions (iii)–(iv),

$$\begin{aligned} \mathbb{E} \left[\sum_{t=1}^T \sqrt{z_t h_t} \right] &\leq \sqrt{\rho} \cdot \varrho_0(\pi_T^\star), \\ \mathbb{E} \left[\sum_{t=1}^T u_t h_t \right] &\leq \rho \cdot \varrho_0(\pi_T^\star). \end{aligned}$$

Furthermore, Lemma 21 of [Kuroki et al. \[2024\]](#) gives the lower bound

$$R_T \geq \frac{\Delta_{\min}}{2} \mathbb{E}[\varrho_0(\pi^*)] - 2C.$$

Balancing both bounds using any $\lambda \in (0, 1]$, and applying the inequalities $ax^2 - bx^3 \leq \frac{4a^3}{27b^2}$ and $ax - bx^2 \leq \frac{a^2}{4b}$ (for $a \geq 0, b > 0$), we obtain after simplification:

$$\begin{aligned} R_T &\lesssim \frac{(1+\lambda)^3}{\lambda^2} \cdot \frac{\rho \log(\varepsilon T)}{\Delta_{\min}^2} + \frac{(1+\lambda)^2}{\lambda} \cdot \frac{\rho \log(\varepsilon T)}{\Delta_{\min}} \\ &\quad + \left(\frac{\sqrt{z_{\max} h_1}}{\varepsilon} \right)^{2/3} + \sqrt{\frac{u_{\max} h_1}{\varepsilon}} + \kappa + 2\lambda C. \end{aligned}$$

Choosing $\lambda = \Theta\left(\left(\frac{\rho \log(\varepsilon T)}{C}\right)^{1/3}\right)$ and setting $\varepsilon = 1/(\rho^2/\Delta_{\min}^3 + C\rho/\Delta_{\min}) \leq 1/T$ gives, for $T \geq \tau := \frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}}$,

$$\begin{aligned} R_T &\lesssim \frac{\rho}{\Delta_{\min}^2} \log_+(T\Delta_{\min}^3) + \left(\frac{C^2 \rho}{\Delta_{\min}^2} \log_+ \left(\frac{T\Delta_{\min}}{C} \right) \right)^{1/3} \\ &\quad + \left((z_{\max} h_1)^{1/3} + \sqrt{u_{\max} h_1} \right) \left(\frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}} \right)^{2/3} + \kappa, \end{aligned}$$

which concludes the proof. \square

C PROOF OF THEOREM 1

We build on Lemma 1, presented and proved in Appendix B, to prove Theorem 1 by verifying that Algorithm 1 satisfies conditions (i)–(iv) of the lemma. We recall the theorem below, before presenting its proof.

Theorem 1. *In the adversarial regime, the regret of Algorithm 1 satisfies*

$$\begin{aligned} R_T &\lesssim \left(\frac{c K d^2 \log K}{\lambda_{\min}^2} \right)^{1/3} T^{2/3} \\ &\quad + \sqrt{\frac{\max(c, 1) d \log K \cdot T}{\lambda_{\min}}} + \kappa \end{aligned}$$

with

$$\begin{aligned} \kappa &= \sqrt{\frac{c K d^2 \log K}{\lambda_{\min}^2}} + \frac{\max(c, 1) d \log K}{\lambda_{\min}} \\ &\quad + \frac{\max(c, 1) K \log K}{\lambda_{\min}^2} + \frac{32 K d \sqrt{c}}{(1-\alpha)^2 \sqrt{\beta_1} \min(1, \lambda_{\min})}. \end{aligned}$$

while in the corrupted stochastic regime with corruption level C it satisfies

$$\begin{aligned} R_T &\lesssim \frac{d \sqrt{\max(c, 1) K \log K}}{\lambda_{\min} \Delta_{\min}^2} \cdot \log(T\Delta_{\min}^3) \\ &\quad + \left(\frac{C^2 d \sqrt{\max(c, 1) K \log K}}{\lambda_{\min} \Delta_{\min}^2} \cdot \log \left(\frac{T\Delta_{\min}}{C} \right) \right)^{1/3} \\ &\quad + \kappa + \kappa', \text{ where we further define} \\ \kappa' &= \left(\left(\frac{c K d^2 \log K}{\lambda_{\min}^2} \right)^{1/3} + \sqrt{\frac{\max(c, 1) d \log K}{\lambda_{\min}}} \right) \\ &\quad \times \left(\frac{1}{\Delta_{\min}^3} + \frac{C}{\Delta_{\min}} \right)^{2/3} \end{aligned}$$

Proof. #DB: Tldr: all this part is shit because it mixes using fixed context and current context, with ambiguous notation that doesn't allow to understand in which case we are. We verify that Algorithm 1 satisfies the four conditions of Lemma 1.

Throughout this proof, we work with the *exact* loss estimates $\hat{\theta}_{t,a}$ defined in Eq. (3), rather than their MGR approximations $\tilde{\theta}_{t,a}$ used in the algorithmic description. This distinction is only technical and does not affect the regret order, since Lemma 9 guarantees that the cumulative bias introduced by the MGR approximation remains uniformly bounded.

Condition (i). By definition of the importance-sampled loss (Eq. (3)), for any $a \in [K]$ we have

$$|\hat{\ell}_{t,a} \eta_t| \leq \frac{\ell_{t,a} \eta_t}{p_t \lambda_{\min}} \leq \frac{1}{u_t \lambda_{\min}} \leq \frac{1-\alpha}{8} \cdot \frac{1}{\min(q_{t,a_t^*}, 1-q_{t,a_t^*})^{1-\alpha}}.$$

Hence, the scaled losses $\hat{\ell}_t \eta_t$ satisfy the condition of Lemma 3, presented in Appendix B, which provides an upper bound on the penalty term $\langle q_t - q_{t+1}, \hat{\ell}_t(x) \rangle - D_t(q_{t+1}, q_t)$ appearing in the standard FTRL regret decomposition.

#DB: Imo, has to be re-written from the start with the ghost sample x and with the estimates $\hat{\theta}_{t,a}$ actually used by FTRL. Since the regret is defined by

$$R_T = \mathbb{E} \left[\sum_{t=1}^T (\langle X_t, \theta_{t,A_t} \rangle - \langle X_t, \theta_{t,\pi^*(X_t)} \rangle) + cK \sum_{t=1}^T p_t \right],$$

and $\hat{\theta}_{t,a}$ is an unbiased estimator of $\theta_{t,a}$, we can equivalently write

$$R_T = \mathbb{E} \left[\sum_{t=1}^T (\langle X_t, \hat{\theta}_{t,A_t} \rangle - \langle X_t, \hat{\theta}_{t,\pi^*(X_t)} \rangle) + cK \sum_{t=1}^T p_t \right].$$

#DB: This is wrong without using the ghost sample technique, cause the estimates depend on X_t (with some X_0 , it's OK).#

Fix any context $x \in \mathbb{R}^d$. Applying Lemma 4, we obtain:

$$\begin{aligned} \sum_{t=1}^T (\langle x, \hat{\theta}_{t,A_t} \rangle - \langle x, \hat{\theta}_{t,\pi^*(x)} \rangle) &\leq \underbrace{\sum_{t=1}^T (\psi_t(q_{t+1}(\cdot|x)) - \psi_{t+1}(q_{t+1}(\cdot|x)))}_{\text{stability}} \\ &\quad + \underbrace{\sum_{t=1}^T (\langle q_t(\cdot|x) - q_{t+1}(\cdot|x), \hat{\ell}_t(x) \rangle - D_t(q_{t+1}(\cdot|x), q_t(\cdot|x))) + A + \bar{\beta}h}_{\text{penalty}}, \end{aligned}$$

where $A = \psi_{T+1}(\pi^*(\cdot|x)) - \psi_1(q_1(\cdot|x)) \leq \beta_1 \log K$ is independent of T and will be ignored in the sequel (together with $\bar{\beta}h$).

Bounding the stability term. By the definition of ψ_t , we have

$$\sum_{t=1}^T (\psi_t(q_{t+1}(\cdot|x)) - \psi_{t+1}(q_{t+1}(\cdot|x))) \leq \sum_{t=1}^T \left(\frac{1}{\eta_{t+1}} - \frac{1}{\eta_t} \right) h_{t+1}.$$

#DB: Now here this is were messing up with the notation is confusing: above check which components should depend on x .# Reindexing $t \mapsto t-1$ yields the equivalent form

$$\sum_{t=1}^T \left(\frac{1}{\eta_t} - \frac{1}{\eta_{t-1}} \right) h_t.$$

Bounding the penalty term. Using Lemma 3 together with Lemma 2, we have

$$\begin{aligned}
 \sum_{t=1}^T (\langle q_t - q_{t+1}, \hat{\ell}_t(x) \rangle - D_t(q_{t+1}, q_t)) &= \sum_{t=1}^T \frac{1}{\eta_t} \left(\langle q_t - q_{t+1}, \hat{\ell}_t(x) \eta_t \rangle - D_t(q_{t+1}, q_t) \right) \\
 &\leq \sum_{t=1}^T \frac{4\eta_t}{1-\alpha} \left(q_{t,a_t^*}^{2-\alpha} \hat{\ell}_{t,a_t^*}^2 + \sum_{a \neq a_t^*} q_{t,a}^{2-\alpha} \hat{\ell}_{t,a}^2 \right) \\
 &\leq \sum_{t=1}^T \frac{4d^2\eta_t}{p_t(1-\alpha)\lambda_{\min}^2} \left(q_{t,a_t^*}^{2-\alpha} + \sum_{a \neq a_t^*} q_{t,a}^{2-\alpha} \right).
 \end{aligned} \tag{12}$$

#DB: Isn't there some $\min(q, 1-q)$ at some point in lemma 3? Otherwise why split between a_t^* and $a \neq a_t^*$. Not consistent#

#DB: More important: lemma 2 bounds $\mathbb{E}[\hat{\ell}_{t,a}^2]$ while here you should have a bound on $\mathbb{E}[q_{t,a}^{2-\alpha} \cdot \hat{\ell}_{t,a}^2]$. Transposing previous works we would work on $\mathbb{E}[\sum_{a=1}^K q_{t,a}^{2-\alpha} \cdot \hat{\ell}_{t,a}^2]$ directly.#

Taking expectations over the random context X_t , we obtain

$$\begin{aligned}
 \mathbb{E}[R_T] &\leq \mathbb{E} \left[\sum_{t=1}^T \left(\frac{1}{\eta_t} - \frac{1}{\eta_{t-1}} \right) h_t \right. \\
 &\quad \left. + \sum_{t=1}^T \frac{4d^2\eta_t}{p_t(1-\alpha)\lambda_{\min}^2} \left(q_{t,a_t^*}^{2-\alpha} + \sum_{a \neq a_t^*} q_{t,a}^{2-\alpha} \right) + cK \sum_{t=1}^T p_t \right],
 \end{aligned}$$

which matches the required structure of condition (i).

#DB: It seems we should not be doing that but instead use a ghost sample to define the regret as in Tsuchiya and Ito [2024], Kuroki et al. [2024].#

Condition (ii). Condition (ii) follows directly from Lemma 6, presented and proved in Appendix B, which guarantees that

$$\mathbb{E}[h_{t+1} \mid \mathcal{H}_t] \leq 2 \mathbb{E}[h_t \mid \mathcal{H}_{t-1}], \quad \forall t \geq 1.$$

Conditions (iii) and (iv). Lemma 13 of Tsuchiya and Ito [2024] provides an upper bound on the entropy term,

$$h_t \leq \frac{1}{\alpha} (K-1)^{1-\alpha} (1 - q_{t,a_t^*})^\alpha,$$

where $a_t^* := \arg \max_{a \in [K]} \langle X_t, \theta_{t,a} \rangle$ denotes the optimal arm for context X_t . Moreover, using the definitions of z_t and u_t , we obtain:

$$\begin{aligned}
 z_t &= \frac{4cKd^2}{(1-\alpha)\lambda_{\min}^2} \left(\sum_{a \neq a_t^*} q_{t,a}^{2-\alpha} + (\min(q_{t,a_t^*}, 1 - q_{t,a_t^*}))^{2-\alpha} \right) \\
 &\leq \frac{8cKd^2}{(1-\alpha)\lambda_{\min}^2} (1 - q_{t,a_t^*})^{2-\alpha}.
 \end{aligned}$$

Combining the bounds on h_t and z_t yields

$$\begin{aligned}
 z_t h_t &\leq \frac{8cKd^2(K-1)^{1-\alpha}}{\alpha(1-\alpha)\lambda_{\min}^2} (1 - q_{t,a_t^*})^2, \\
 u_t h_t &\leq \frac{8d \max(c, 1)}{(1-\alpha)\alpha} (K-1)^{1-\alpha} (1 - q_{t,a_t^*}).
 \end{aligned}$$

Hence, both conditions are satisfied with

$$\begin{aligned}
 \sqrt{z_t h_t} &\leq \sqrt{\rho} (1 - q_{t,a_t^*}), \\
 u_t h_t &\leq \rho (1 - q_{t,a_t^*}),
 \end{aligned}$$

where

$$\rho := \frac{d}{\lambda_{\min}} \max \left(\sqrt{\frac{8cK(K-1)^{1-\alpha}}{\alpha(1-\alpha)}}, \frac{8 \max(c, 1)}{(1-\alpha)\alpha} (K-1)^{1-\alpha} \right).$$

Conclusion. Having verified conditions (i)–(iv), we can invoke Lemma 1 to conclude that Algorithm 1 enjoys a Best-of-Both-Worlds (BoBW) regret guarantee. To make the constants explicit, note that

$$h_{\max} \leq \frac{K^{1-\alpha}}{\alpha}, \quad z_{\max} = \mathcal{O}\left(\frac{cKd^2}{(1-\alpha)\lambda_{\min}^2}\right), \quad u_{\max} = \mathcal{O}\left(\frac{d \max(c, 1)}{1-\alpha}\right).$$

Plugging these into Lemma 1 gives

$$\begin{aligned} \text{Adversarial regime: } R_T &= \mathcal{O}\left(\left(\frac{cKd^2}{\lambda_{\min}^2}\right)^{1/3} T^{2/3} + \sqrt{\frac{d \max(c, 1)T}{\lambda_{\min}}}\right), \\ \text{Corrupted stochastic regime: } R_T &= \mathcal{O}\left(\frac{d\sqrt{\max(c, 1)K}}{\lambda_{\min}\Delta_{\min}^2} \log(T\Delta_{\min}^3) + \left(\frac{C^2 d\sqrt{\max(c, 1)K}}{\lambda_{\min}\Delta_{\min}^2} \log\frac{T\Delta_{\min}}{C}\right)^{1/3}\right). \end{aligned}$$

This completes the proof of Theorem 1. \square

D TECHNICAL LEMMAS

#DB: Below we have a problem, $\lambda_{\min}(\Sigma_{t,a}) \geq p_t \lambda_{\min}$ has no reason to hold and this should be fixed.#

#DB: Additionally, we haven't sorted the dimension issue yet. What should be the right dependence in d ? It's not so clear, but it should appear in the squared loss if we follow previous works: there, d come from using a trace bound instead of our inequalities below. Some of our steps are probably wrong.#

#DB: The following result should bound $\mathbb{E}[\sum_a q_{t,a}^{2-\alpha} \hat{\ell}_{t,a}^2]$ or whatever is needed in the part of the proof where it's used.#

Lemma 2. Let $X_t \in \mathbb{R}^d$ be a random context and fix any arm $a \in [K]$. Under the assumptions of Section 2, we have $\|X_t\|_2 \leq 1$ almost surely, and the loss function satisfies $-1 \leq \ell_t(X_t, a) \leq 1$.

We also recall that $\Sigma_{t,a}$ is a positive definite matrix such that

$$\lambda_{\min}(\Sigma_{t,a}) \geq p_t \lambda_{\min},$$

#DB: We should probably adopt a notation $p_t(X)$, a lot of confusion comes from using p_t for many different things. The above line is clearly wrong (if p_t is $p_t(X_t)$). It is true with $p_{t,\min} := \min_{x \in \mathcal{X}} p_t(x)$ though, and could be replaced by $\frac{\gamma}{K}$ if we add a forced exploration γ .# and that the importance-weighted estimator is given by

$$\hat{\theta}_{t,a} := \Sigma_{t,a}^{-1} X_t \ell_t(X_t, a) \mathbf{1}\{a \in O_t\},$$

where $\mathbb{P}(a \in O_t) = p_t$. Then,

$$\mathbb{E}\left[\langle X_t, \hat{\theta}_{t,a} \rangle^2\right] \leq \frac{1}{\lambda_{\min}^2 p_t}.$$

Proof. We know that the smallest eigenvalue of $\Sigma_{t,a}$ is $\geq p_t \lambda_{\min}$.

$$\|\Sigma_{t,a}^{-1}\|_2 \leq \frac{1}{\lambda_{\min}(\Sigma_{t,a})} \leq \frac{1}{p_t \lambda_{\min}}.$$

Therefore,

$$\|\hat{\theta}_{t,a}\|_2 = \|\Sigma_{t,a}^{-1} X_t \ell_t(X_t, a) \mathbf{1}\{a \in O_t\}\|_2 \leq \frac{\|X_t\|_2}{p_t \lambda_{\min}} \mathbf{1}\{a \in O_t\}.$$

By the Cauchy–Schwarz inequality,

$$\langle X_t, \hat{\theta}_{t,a} \rangle^2 \leq \|X_t\|_2^2 \|\hat{\theta}_{t,a}\|_2^2 \leq \frac{\|X_t\|_2^4}{p_t^2 \lambda_{\min}^2} \mathbf{1}\{a \in O_t\}.$$

Taking expectations and using $\mathbb{E}[\mathbf{1}\{a \in O_t\}] = p_t$, we obtain

$$\mathbb{E}[\langle X_t, \hat{\theta}_{t,a} \rangle^2] \leq \frac{\mathbb{E}[\|X_t\|_2^4]}{p_t \lambda_{\min}^2}.$$

Since $\|X_t\|_2 \leq \sqrt{d}$ almost surely, it follows that $\mathbb{E}[\|X_t\|_2^4] \leq 1$, and hence

$$\mathbb{E}[\langle X_t, \hat{\theta}_{t,a} \rangle^2] \leq \frac{1}{\lambda_{\min}^2 p_t}.$$

□

#DB: (minor for now) Write a few sentences between lemmas. Also, here we just do an unstructured list, but actually the above lemma is crucial, while the following is just a restatement of an existing result. Might be nice to structure this. #

Lemma 3 (Lemma 14 in [Tsuchiya and Ito \[2024\]](#)). *Let $q \in \mathcal{P}_K$ and let $\bar{I} \in \arg \max_{i \in [K]} q_i$. Let $l \in \mathbb{R}^K$ be such that, for all $i \in [K]$,*

$$|l_i| \leq \frac{1-\alpha}{4} \cdot \frac{1}{\min(q_{\bar{I}}, 1-q_{\bar{I}})^{1-\alpha}}.$$

Then, the following bound holds:

$$\max_{p \in \mathcal{P}_K} \{ \langle l, q - p \rangle - D_{-H_\alpha}(p, q) \} \leq \frac{4}{1-\alpha} \left(\sum_{i \neq \bar{I}} q_i^{2-\alpha} l_i^2 + \min(q_{\bar{I}}, 1-q_{\bar{I}})^{2-\alpha} l_{\bar{I}}^2 \right)$$

Lemma 4. *Let $x \in \mathbb{R}^d$ be any fixed context. For each $t \geq 1$, let $q_t(\cdot|x) \in \Delta_K$ be the distribution used to sample A_t given x , and let $\pi^*(\cdot|x) \in \Delta_K$ be any comparator policy. Let $(\psi_t)_{t \geq 1}$ be a sequence of σ -strongly convex regularizers on Δ_K , and let $D_t(\cdot, \cdot)$ denote the Bregman divergence induced by ψ_t . Denoting $\hat{\ell}_t(x) \in \mathbb{R}^K$ for the vector of estimated losses at context x , with $[\hat{\ell}_t(x)]_a := \langle x, \hat{\theta}_{t,a} \rangle$. Then*

$$\begin{aligned} \mathbb{E} \left[\sum_{t=1}^T \left(\langle x, \hat{\theta}_{t,A_t} \rangle - \langle x, \hat{\theta}_{t,\pi^*(x)} \rangle \right) \right] &\leq \mathbb{E} \left[\sum_{t=1}^T (\psi_t(q_{t+1}(\cdot|x)) - \psi_{t+1}(q_{t+1}(\cdot|x))) \right] \\ &\quad + \mathbb{E} \left[\sum_{t=1}^T \left(\langle q_t(\cdot|x) - q_{t+1}(\cdot|x), \hat{\ell}_t(x) \rangle - D_t(q_{t+1}(\cdot|x), q_t(\cdot|x)) \right) \right] \\ &\quad + \mathbb{E}[\psi_{T+1}(\pi^*(\cdot|x)) - \psi_1(q_1(\cdot|x))] + \bar{\beta} \bar{h}, \end{aligned}$$

where $\bar{h} := \max_{p \in \Delta_K} H_{\bar{\alpha}}(p)$ and $\bar{\beta} \geq 0$ is the coefficient that upper-bounds the change of regularizer in our setting.

#DB: This lemma above uses a fixed context, so is suited for a proof that uses the ghost sample technique. It cannot work for the proof as it is right now.#

#DB: Has to be modified if forced exploration is introduced. See (44) in [Kuroki et al. \[2024\]](#) for their result.#

Proof. Conditionally on x , $A_t \sim q_t(\cdot|x)$, hence

$$\mathbb{E}[\langle x, \hat{\theta}_{t,A_t} \rangle | x] = \sum_{a \in [K]} q_t(a|x) \langle x, \hat{\theta}_{t,a} \rangle.$$

Therefore,

$$\begin{aligned} \mathbb{E} \left[\sum_{t=1}^T \left(\langle x, \hat{\theta}_{t,A_t} \rangle - \langle x, \hat{\theta}_{t,\pi^*(x)} \rangle \right) \right] &= \mathbb{E} \left[\sum_{t=1}^T \sum_{a \in [K]} (q_t(a|x) - \pi^*(a|x)) \langle x, \hat{\theta}_{t,a} \rangle \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T \langle q_t(\cdot|x) - \pi^*(\cdot|x), \hat{\ell}_t(x) \rangle \right]. \end{aligned}$$

We now invoke the standard FTRL regret decomposition with time-varying regularizers (see, e.g. Exercise 28.12 in Lattimore and Szepesvári [2020]): **#DB:** I think this works only with the estimate that is actually used in FTRL. So, for us, $\tilde{\theta}_{t,a}$.# for any $q \in \Delta_K$,

$$\begin{aligned} \sum_{t=1}^T \langle q_t(\cdot|x) - q, \hat{\ell}_t(x) \rangle &\leq \psi_{T+1}(q) - \psi_1(q_1) + \sum_{t=1}^T (\psi_t(q_{t+1}(\cdot|x))) - \psi_{t+1}(q_{t+1}(\cdot|x)) \\ &+ \sum_{t=1}^T (\langle q_t(\cdot|x) - q_{t+1}(\cdot|x), \hat{\ell}_t(x) \rangle - D_t(q_{t+1}(\cdot|x), q_t(\cdot|x))). \end{aligned}$$

Choosing $q = \pi^*(\cdot|x)$ and taking expectations yields the claim, with the additional additive term $\bar{\beta} \bar{h}$ accounting for the regularizer variation bound used in our setup. \square

Lemma 5 (Lemma 15 of Tsuchiya and Ito [2024]). *Let $l, L \in \mathbb{R}^K$, let $q, r \in \mathcal{P}_k$ be:*

$$q \in \arg \min_{p \in \mathcal{P}_k} \{ \langle L, p \rangle + \beta(-H_\alpha(p)) + \bar{\beta}(-H_{\bar{\alpha}}(p)) \}$$

$$r \in \arg \min_{p \in \mathcal{P}_k} \{ \langle L + l, p \rangle + \beta'(-H_\alpha(p)) + \bar{\beta}(-H_{\bar{\alpha}}(p)) \}$$

for the Tsallis entropy H_α and $0 < \beta < \beta'$. Suppose also that

$$\|l\|_\infty \leq \max\left(\frac{1 - (\sqrt{2})^{\alpha-1}}{2} q_*^{\alpha-1} \beta, \frac{1 - (\sqrt{2})^{\bar{\alpha}-1}}{2} q_*^{\bar{\alpha}-1} \bar{\beta}\right)$$

$$0 \leq \beta' - \beta \leq \max((1 - (\sqrt{2})^{\alpha-1})\beta, \frac{1 - (\sqrt{2})^{\bar{\alpha}-1}}{\sqrt{2}} q_*^{\bar{\alpha}-\alpha} \bar{\beta})$$

Then it holds that $H_\alpha(r) \leq 2H_\alpha(q)$.

#DB: Check carefully this thing below, cause there might be again the issue of treating some probabilities as fixed while they're not. Note that the ghost sample technique might require that we only have to control the entropy of a fixed context X_0 , which might simplify things.#

Lemma 6. *Algorithm 1 satisfies, for all $t \geq 1$,*

$$\mathbb{E}[h_{t+1} | \mathcal{H}_t] \leq 2 \mathbb{E}[h_t | \mathcal{H}_{t-1}].$$

Proof. We first control the key quantities appearing in Lemma 5. Recall that $\beta_t = 1/\eta_t$, $\gamma_t = \sqrt{z_t/\beta_t} + u_t/\beta_t$, and $h_t = \frac{1}{\alpha} \sum_{i=1}^K (q_{t,i}^\alpha - q_{t,i})$.

Step 1: Bounding $\sqrt{z_t}$ and h_t . By definition of z_t we have

$$\sqrt{z_t} = \sqrt{\frac{4cKd^2}{1-\alpha} \left(\sum_{i \neq I_t} q_{t,i}^{2-\alpha} + q_{t,a_t^*}^{2-\alpha} \right)} \leq \frac{2d\sqrt{Kc}}{\sqrt{1-\alpha}} q_{t,a_t^*}^{1-\frac{\alpha}{2}}.$$

In addition, from the properties of the Tsallis entropy (see, e.g., Lemma 13 of Tsuchiya and Ito [2024]),

$$h_t = \frac{1}{\alpha} \sum_{i=1}^K (q_{t,i}^\alpha - q_{t,i}) \geq \frac{1 - (1/2)^{1-\alpha}}{\alpha} q_{t,a_t^*}^\alpha \geq \frac{1-\alpha}{4\alpha} q_{t,a_t^*}^\alpha.$$

Step 2: Bounding the variation of β_t . From Equation 9,

$$\beta_{t+1} - \beta_t = \frac{2}{h_t} \sqrt{\frac{z_t}{\beta_t}} + \frac{u_t}{h_t \beta_t}.$$

Plugging in the bounds on $\sqrt{z_t}$ and h_t gives

$$\begin{aligned} \beta_{t+1} - \beta_t &\leq \frac{16\alpha d \sqrt{Kc}}{\sqrt{\beta_t}(1-\alpha)^{3/2}} q_{t,a_t^*}^{1-\frac{3\alpha}{2}} + \frac{32\alpha d \max(c,1)}{\sqrt{\beta_t}(1-\alpha)^2 \lambda_{\min}} q_{t,a_t^*}^{1-2\alpha} \\ &\leq \alpha \bar{\beta} q_{t,a_t^*}^{1-\frac{3\alpha}{2}} + \frac{\alpha \bar{\beta}}{\lambda_{\min}} q_{t,a_t^*}^{1-2\alpha} \\ &\leq 2 \frac{(1-\bar{\alpha})}{\min(1, \lambda_{\min})} \bar{\beta} q_{t,a_t^*}^{\bar{\alpha}-\alpha} \leq 2 \frac{1-(\sqrt{2})^{\bar{\alpha}-1}}{\sqrt{2}} \bar{\beta} q_{t,a_t^*}^{\bar{\alpha}-\alpha}. \end{aligned}$$

Hence, $\beta_{t+1} - \beta_t$ satisfies the second condition of Lemma 5.

Step 3: Bounding the loss magnitude. For any fixed context x and arm $i \in [K]$,

$$\begin{aligned} |\hat{\ell}_{t+1,i}(x)| &= |\langle x, \hat{\theta}_{t+1,i} \rangle| \leq \frac{d}{\lambda_{\min} p_t} \leq \frac{d}{\lambda_{\min}} \cdot \frac{\beta_t}{u_t} \\ &= \frac{1-\alpha}{8} \cdot \frac{\beta_t}{q_{t,a_t^*}^{1-\alpha}} \leq \frac{1-(\sqrt{2})^{\alpha-1}}{2} \cdot \frac{\beta_t}{q_{t,a_t^*}^{1-\alpha}}. \end{aligned}$$

This matches the first smoothness condition of Lemma 5.

Step 4: Applying Lemma 5. Since both smoothness conditions are satisfied, the lemma implies

$$H_\alpha(q_{t+1}) \leq 2 H_\alpha(q_t),$$

and therefore $h_{t+1} \leq 2h_t$ whenever the context remains fixed.

Taking conditional expectations and using the stationarity of the context distribution then yields

$$\mathbb{E}[h_{t+1} | \mathcal{H}_t] \leq 2 \mathbb{E}[h_t | \mathcal{H}_{t-1}],$$

which completes the proof. \square

Lemma 7 (Slight adaptation of Theorem 6 of [Tsuchiya and Ito \[2024\]](#)). *For all $\varepsilon \geq 1/T$, it holds that*

$$\begin{aligned} &F(\beta_{1:T}, z_{1:T}, u_{1:T}, h_{0:T-1}) \\ &\lesssim \min \left\{ \left(\sum_{t=1}^T \sqrt{z_t h_t \log(\varepsilon T)} \right)^{2/3}, \right. \\ &\quad \left. \left(\frac{\sqrt{z_{\max} h_{\max}}}{\varepsilon} \right)^{2/3}, \left(\sum_{t=1}^T \sqrt{z_t h_{\max}} \right)^{2/3} \right\} \\ &\quad + \min \left\{ \sqrt{\sum_{t=1}^T u_t h_t \log(\varepsilon T)}, \frac{\sqrt{u_{\max} h_{\max}}}{\varepsilon}, \sum_{t=1}^T u_t h_{\max} \right\} \\ &\quad + \sqrt{\frac{z_{\max}}{\beta_1}} + \frac{u_{\max}}{\beta_1} + \beta_1 h_1 \end{aligned}$$

Proof. This slight adaptation originates from a minor modification of Lemma 4 in [Tsuchiya and Ito \[2024\]](#), where in the first line of equation (24) we instead bound:

$$\begin{aligned} F(\beta_{1:T}, z_{1:T}, u_{1:T}, h_{0:T-1}) &\leq 2\sqrt{\frac{z_1}{\beta_1}} + \frac{u_1}{\beta_1} + \beta_1 h_1 \\ &+ \sum_{t=2}^T \left(2\sqrt{\frac{z_t}{\beta_t}} + \frac{u_t}{\beta_t} + (\beta_t - \beta_{t-1})h_{t-1} \right). \end{aligned}$$

After this adjustment, the remainder of the proof proceeds identically. \square

E MATRIX GEOMETRIC RESAMPLING

Before detailing Algorithm 2, we elaborate on why using the parameter estimates from Eq. (3) would be untractable in practice. To prove this point, we detail the computation of the exact covariance matrix $\Sigma_{t,a}$, which involves evaluating the following conditional expectation:

$$\begin{aligned} \Sigma_{t,a} &= \mathbb{E}_t[\mathbf{1}_{a \in O_t} X_t X_t^\top] = \sum_{X \in \mathcal{X}} \mathbb{P}_{X_t,a}(X_t = X, a \in O_t) XX^\top \\ &= \sum_{X \in \mathcal{X}} \mathbb{P}_{X_t \sim \mathcal{D}}(X_t = X) \underbrace{\mathbb{P}(a \in O_t | X_t = X)}_{p_t(X)} XX^\top. \end{aligned}$$

The challenge lies in evaluating the conditional observation probability $p_t(X)$. Note that in Algorithm 1, p_t was defined unambiguously since it was the observation probability corresponding to the (unique) fixed context X_t , computed after it is revealed. Here, $p_t(X)$ is derived following the same steps, but computed as if context X was observed instead of X_t . Doing so requires performing all computations leading to Eq. (8) separately for each possible context $X \in \mathcal{X}$. This results in a computational complexity proportional to the size of the context space, $|\mathcal{X}|$, which becomes quickly prohibitive when \mathcal{X} is large. In addition, we can note that each individual computation requires solving an optimization problem to obtain the FTRL sampling probability (Eq. (2)).

To circumvent this limitation, we follow [Neu and Bartók \[2013\]](#), [Neu \[2015\]](#), [Kuroki et al. \[2024\]](#) and use Matrix Geometric Resampling (MGR) to efficiently approximate the *inverse* of the matrix $\Sigma_{t,a}$ directly. It doesn't need to compute the FTRL sampling allocation over all possible contexts but only on a carefully chosen number of sampled contexts, and only use matrix products (costing $\mathcal{O}(d^2)$) but no matrix inversion (costing $\mathcal{O}(d^3)$). We recall this procedure in Algorithm 2 below. In the pseudo-code, we denote by $\mathcal{B}(p)$ the Bernoulli distribution with parameter p .

Algorithm 2 Matrix Geometric Resampling (MGR)

Require: Sampler of the context distribution \mathcal{D} , number of iterations M_t

- 1: Initialize $\Sigma_t^+ \leftarrow \frac{1}{2}I$, $A_0 = I$
 - 2: **for** $i = 1$ to M_t **do**
 - 3: Sample $X \sim \mathcal{D}$
 - 4: Compute probability of observation p as in Step 5 of Algorithm 1 if X_t was equal to X .
 - 5: Sample $b \sim \mathcal{B}(p)$
 - 6: Compute $B_i \leftarrow bXX^\top$
 - 7: Compute $A_i \leftarrow A_{i-1}(I - \frac{1}{2}B_i)$
 - 8: Update $\Sigma_t^+ \leftarrow \Sigma_t^+ + \frac{1}{2}A_i$
 - 9: **end for**
 - 10: **return** Σ_t^+
-

We now introduce the technical results related to the cost and approximation guarantees of the MGR procedure, which will be used in the regret analysis (see the proof sketch in Section 4).

Lemma 8 (Adapted from Lemma 9 of Kuroki et al. [2024]). Let $\widehat{\theta}_{t,a} = \Sigma_{t,a}^{-1} X_t l_t(X_t, A_t) \mathbb{I}\{a \in O_t\}$ and let $\widetilde{\theta}_{t,a} = \Sigma_{t,a}^+ X_t l_t(X_t, A_t) \mathbb{I}\{a \in O_t\}$, where $\Sigma_{t,a}^+$ is obtained via the Matrix Geometric Resampling (MGR) procedure in Algorithm 2 with the number of iterations M_t tuned as in Eq. (6). Then, for any arm $a \in [K]$ and round $t \geq 1$, it holds that

$$\mathbb{E} \left[\langle X_t, \widetilde{\theta}_{t,a} - \widehat{\theta}_{t,a} \rangle \mid \mathcal{H}_{t-1} \right] \leq \exp \left(-\frac{p_t \lambda_{\min}}{2K} M_t \right).$$

#DB: Again, could only be used by plugging $p_{t,\min}$ or forced exploration instead of p_t .#

Proof. Let $\|\cdot\|_{\text{op}}$ denote the operator norm. Denote by $\widehat{\Sigma}_{t,a}^+$ the random matrix output by the MGR procedure in Algorithm 2. Under independence assumptions of the geometric resampling steps, we have

$$\mathbb{E} \left[\prod_{j=1}^i \left(I - \frac{1}{2} B_j \right) \right] = \left(I - \frac{1}{2} \Sigma_{t,a} \right)^i,$$

and consequently,

$$\mathbb{E} \left[\widehat{\Sigma}_{t,a}^+ \right] = \frac{1}{2} \sum_{i=0}^{M_t} \left(I - \frac{1}{2} \Sigma_{t,a} \right)^i = \Sigma_{t,a}^{-1} - \left(I - \frac{1}{2} \Sigma_{t,a} \right)^{M_t} \Sigma_{t,a}^{-1}.$$

Using this, we compute the expectation of the biased estimator:

$$\begin{aligned} \mathbb{E}[\widetilde{\theta}_{t,a}] &= \mathbb{E}[\widehat{\Sigma}_{t,a}^+ X_t l_t(X_t, a) \mathbb{I}\{A_t = a\}] \\ &= \mathbb{E}[\widehat{\Sigma}_{t,a}^+] \cdot \mathbb{E}[X_t \langle X_t, \theta_{t,a} \rangle \mathbb{I}\{A_t = a\}] \\ &= \mathbb{E}[\widehat{\Sigma}_{t,a}^+] \cdot \mathbb{E}[X_t X_t^\top \mathbb{I}\{A_t = a\}] \cdot \theta_{t,a} \\ &= \left(\Sigma_{t,a}^{-1} - \left(I - \frac{1}{2} \Sigma_{t,a} \right)^{M_t} \Sigma_{t,a}^{-1} \right) \cdot \Sigma_{t,a} \cdot \theta_{t,a} \\ &= \theta_{t,a} - \left(I - \frac{1}{2} \Sigma_{t,a} \right)^{M_t} \theta_{t,a}. \end{aligned}$$

Hence, the bias is given by:

$$\mathbb{E}[\widetilde{\theta}_{t,a} - \widehat{\theta}_{t,a}] = - \left(I - \frac{1}{2} \Sigma_{t,a} \right)^{M_t} \theta_{t,a}.$$

We then bound the inner product as:

$$\begin{aligned} \left| \mathbb{E} \left[\langle X_t, \widetilde{\theta}_{t,a} - \widehat{\theta}_{t,a} \rangle \mid \mathcal{H}_{t-1} \right] \right| &\leq \|X_t\|_2 \cdot \|\theta_{t,a}\|_2 \cdot \left\| \left(I - \frac{1}{2} \Sigma_{t,a} \right)^{M_t} \right\|_{\text{op}} \leq \left\| \left(I - \frac{1}{2} \Sigma_{t,a} \right)^{M_t} \right\|_{\text{op}} \\ &\leq \left(1 - \frac{p_t \lambda_{\min}}{2K} \right)^{M_t} \leq \exp \left(-\frac{p_t \lambda_{\min}}{2K} M_t \right), \end{aligned}$$

where we used $\|X_t\|_2 \leq 1$, $\|\theta_{t,a}\|_2 \leq 1$, and the bound $\Sigma_{t,a} \succeq \frac{p_t \lambda_{\min}}{K} I$ in the third inequality (since each arm is observed with probability p_t). #DB: Why do we have a K^{-1} factor here?# \square

Lemma 9. The cumulative bias introduced by the MGR approximation is uniformly bounded as

$$\sum_{t=1}^T \max_{a \in [K]} \left| \mathbb{E}[\langle X_t, \widetilde{\theta}_{t,a} - \widehat{\theta}_{t,a} \rangle] \right| \leq \frac{\pi^2}{6}.$$

Proof. From Lemma 8 and the definition $M_t = \lceil \frac{4K}{p_t \lambda_{\min}} \log t \rceil$, we obtain, conditionally on \mathcal{H}_{t-1} ,

$$\left| \mathbb{E}[\langle X_t, \widetilde{\theta}_{t,a} - \widehat{\theta}_{t,a} \rangle \mid \mathcal{H}_{t-1}] \right| \leq \exp \left(-\frac{p_t \lambda_{\min}}{2K} M_t \right) \leq \frac{1}{t^2}.$$

Taking total expectation and maximizing over $a \in [K]$ yields

$$\max_{a \in [K]} \left| \mathbb{E}[\langle X_t, \tilde{\theta}_{t,a} - \hat{\theta}_{t,a} \rangle] \right| \leq \frac{1}{t^2}.$$

We finally obtain the result by summing over t . \square

F TIME AND SPACE COMPLEXITY OF ALGORITHM 1

#DB: Propagate the computations below in the main text. #

In this section, we detail the computation of the memory requirement and computation time of Algorithm 1, presented at the end of Section 3 of the paper.

At each round t , the algorithm stores the tuple (X_t, A_t, p_t, q_t) , which is of negligible size $\mathcal{O}(d + K)$, together with the parameter estimates $\hat{\theta}_{t,a}$ for all $a \in [K]$ and $t \leq T$, which must be kept across rounds to enable information reuse. This requires a total of $\mathcal{O}(dKT)$ memory. In addition, at each round t , computing the MGR approximation requires storing $\Sigma_t^+ \in \mathbb{R}^{d \times d}$, which leads to an additional temporary cost of $\mathcal{O}(d^2)$ during the computation of that round. Therefore, with Equation (13), the total space complexity is

$$\mathcal{O}(dKT + d^2).$$

Per-round computational cost. Each round involves two main computational steps: (i) solving the FTRL update via convex optimization, and (ii) performing Matrix Geometric Resampling (MGR).

FTRL update. In practice, we solve the FTRL objective up to precision $\varepsilon_t = \mathcal{O}(1/t^2)$ so that the cumulative optimization error remains finite. Because of that, we ignored this term in the regret bound of Theorem 1, which assumes that the computation of the sampling distribution is exact.

Using projected gradient descent, the number of iterations required at round t is $\mathcal{O}(\log t)$, and each iteration costs $\mathcal{O}(d \log d)$. Hence, the total cost over T rounds, that we denote by $\text{Comp}_T^{\text{FTRL}}$, satisfies

$$\mathcal{O}\left(\sum_{t=1}^T d \log t \log d\right) = \mathcal{O}(Td \log T \log d).$$

Matrix Geometric Resampling. We recall that M_t denotes the number of resampling steps performed at round t . Let us assume that $1/p_t = \mathcal{O}(t)$, which essentially corresponds to assuming that the logarithmic regret bound of Theorem 1 is also a lower bound, which is reasonable from an information-theoretic perspective. Then, since by Eq. (6) we defined M_t such that

$$M_t = \mathcal{O}\left(\frac{Kt \log t}{\lambda_{\min}}\right), \quad (13)$$

we can define the total computational cost of the MGR procedure at round t as

$$\Gamma_t^{\text{MGR}} \lesssim d^2 M_t \lesssim \frac{Kd^2 t \log t}{\lambda_{\min}}.$$

Summing over all rounds up to T yields a total computation time $\text{Comp}_T^{\text{MGR}}$ satisfying

$$\text{Comp}_T^{\text{MGR}} = \sum_{t=1}^T \Gamma_t^{\text{MGR}} \lesssim \frac{Kd^2}{\lambda_{\min}} \sum_{t=1}^T t \log t \lesssim \frac{Kd^2 T^2 \log T}{\lambda_{\min}}.$$

Thus, since the MGR procedure has to be fully rerun at each iteration, its total computational cost scales quadratically in T .

Overall complexity. Combining the two components of the algorithm, we obtain a total computational cost of order $\text{Comp}_T^{\text{total}} = \text{Comp}_T^{\text{FTRL}} + \text{Comp}_T^{\text{MGR}}$, satisfying

$$\text{Comp}_T^{\text{total}} \lesssim Td \log T \log d + \frac{Kd^2 T^2 \log T}{\lambda_{\min}}.$$

Treating λ_{\min}^{-1} as a numerical constant, and remarking that the second term (MGR steps) dominates, we obtain that the overall running time of the algorithm scales as

$$\text{Comp}_T^{\text{total}} \lesssim Kd^2 T^2 \log T.$$